

第1章 PyTorch程序的基本结构

主要内容

PyTorch的发展历史

参考 <https://mp.weixin.qq.com/s/JrutTVvFtx3xZoagy661LQ>

- 相关的人: Soumith Chintala

Torch7

PyTorch的启动

加入Apache基金会

###

下面是一个非常简单的PyTorch训练代码

```
import os
import time

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.autograd import Variable

from torch.utils.data import DataLoader
from torchvision import datasets, transforms

from collections import OrderedDict
import torch.utils.model_zoo as model_zoo
from torchvision import models

def get_dataset(batch_size, data_root='/tmp/public_dataset/pytorch', train=True, val=True, *
               data_root = os.path.expanduser(os.path.join(data_root, 'mnist-data'))

ds = []
if train:
    train_loader = torch.utils.data.DataLoader(
        datasets.MNIST(root=data_root, train=True, download=True,
                        transform=transforms.Compose([
```

```

        transforms.Resize((224, 224)),
        transforms.Grayscale(3),
        transforms.ToTensor(),
        transforms.Normalize((0.1307,), (0.3081,))
    ])),
    batch_size=batch_size, shuffle=True, **kwargs)
ds.append(train_loader)
if val:
    test_loader = torch.utils.data.DataLoader(
        datasets.MNIST(root=data_root, train=False, download=True,
            transform=transforms.Compose([
                transforms.Resize((224, 224)),
                transforms.Grayscale(3),
                transforms.ToTensor(),
                transforms.Normalize((0.1307,), (0.3081,))
            ])),
        batch_size=batch_size, shuffle=True, **kwargs)
    ds.append(test_loader)
ds = ds[0] if len(ds) == 1 else ds
return ds

epochs = 10
test_interval = 1
data_root = 'data'

use_cuda = torch.cuda.is_available()

# data loader
train_loader, test_loader = get_dataset(batch_size=200, data_root='./data', num_workers=1)

# model
model = models.resnet18(pretrained=True)
in_features = model.fc.in_features
model.fc = nn.Linear(in_features, 10)
if use_cuda:
    model.cuda()

# optimizer
optimizer = optim.SGD(model.parameters(), lr=0.01, weight_decay=0.0001, momentum=0.9)

t_begin = time.time()

for epoch in range(epochs):
    model.train()

```

```

total = 0
for batch_idx, (data, target) in enumerate(train_loader):
    indx_target = target.clone()
    if use_cuda:
        data, target = data.cuda(), target.cuda()

    optimizer.zero_grad()
    output = model(data)
    loss = F.cross_entropy(output, target)
    loss.backward()
    optimizer.step()

    total += len(data)
    elapse_time = time.time() - t_begin
    t_begin = elapse_time
    print("samples {}, time {}s".format(total, int(elapse_time)))

if epoch % test_interval == 0:
    model.eval()
    test_loss = 0
    correct = 0
    for data, target in test_loader:
        indx_target = target.clone()
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        output = model(data)
        test_loss += F.cross_entropy(output, target).data
        pred = output.data.max(1)[1] # get the index of the max log-probability
        correct += pred.cpu().eq(indx_target).sum()

    test_loss = test_loss / len(test_loader) # average over number of mini-batch
    acc = 100. * correct / len(test_loader.dataset)
    print('Test set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)'.format(
        test_loss, correct, len(test_loader.dataset), acc))

```

从这段代码可以看到，一般模型训练的代码包括几个部分： * 数据集的处理和加载
 * 神经网络结构的构建、初始化 * 优化器的配置 * 损失函数的选择，见line
 79，这里用的是交叉熵 * 迭代训练并定期在验证集上测试验证其准确率 *
 保存合适的模型文件，这里没有做这一步

PyTorch的源代码结构

PyTorch的整体架构

PyTorch的源代码结构

```
pytorch
|--- android      # PyTorch for Android
|--- aten         # C++ Tensor
|--- benchmarks   # PyTorch Benchmarking
|--- binaries     #
|--- c10          # Tensor
|--- caffe2       # Caffe2
|--- cmake        # PyTorch
|--- docs         # PyTorch      Python C++
|--- ios          # PyTorch for iOS
|--- modules      #
|--- mypy_plugins #
|--- scripts      #
|--- submodules   #
|--- test         #
|--- third_party  #
|--- tools        #
|--- torch        # PyTorch Python
|--- torchgen     #

torch
|--- csrc         # Torch C++
    |--- module.cpp # Torch C++
```

C10

C10，来自于Caffe Tensor Library的缩写。这里存放的都是最基础的Tensor库的代码，可以运行在服务端和移动端。

C10目前最具代表性的一个class就是TensorImpl了，它实现了Tensor的最基础框架。继承者和使用者有：

```
Variable Variable::Impl
SparseTensorImpl
detail::make_tensor<TensorImpl>(storage_impl, CUDATensorId(), false)
Tensor(c10::intrusive_ptr<TensorImpl, UndefinedTensorImpl> tensor_impl)
c10::make_intrusive<at::TensorImpl, at::UndefinedTensorImpl>
```

值得一提的是，C10中还使用/修改了来自llvm的SmallVector，在vector元素比较少的时候用以代替std::vector，月

ATen

ATen, 来自于 A TENSOR library for C++11的缩写; PyTorch的C++ tensor library。ATen部分有大量的代码是来声明和定义Tensor运算相关的逻辑的, 除此之外, PyTorch还使用了aten/src/ATen

Caffe2

为了复用, 2018年4月Facebook宣布将Caffe2的仓库合并到了PyTorch的仓库, 从用户层面来复用包含了代码、CI、部分37m-x86_64-linux-gnu.so (caffe2 CPU Python 绑定)、caffe2_pybind11_state_gpu.cpython-37m-x86_64-linux-gnu.so (caffe2 CUDA Python 绑定), 基本上来自旧的caffe2项目)

Torch

Torch, 部分代码仍然在使用以前的快要进入历史博物馆的Torch开源项目, 比如具有下面这些文件名格式的文件:

```
TH* = Torch
THC* = Torch Cuda
THCS* = Torch Cuda Sparse (now defunct)
THCUNN* = Torch CUda Neural Network (see cunn)
THD* = Torch Distributed
THNN* = Torch Neural Network
THS* = Torch Sparse (now defunct)
THP* = Torch Python
```

PyTorch会使用tools/setup_helpers/generate_code.py来动态生成Torch层面相关的一些代码, 这部分动态生成的代码

参考

- PyTorch ATen代码的动态生成 <https://zhuanlan.zhihu.com/p/55966063>
- Pytorch1.3源码解析-第一篇 <https://www.cnblogs.com/jeshy/p/11751253.html>

第四章 PyTorch的编译

主要内容

- PyTorch的编译过程
- setup.py的结构
- 代码生成过程
- 生成的二进制包

环境准备

大多数情况下我们只需要安装PyTorch的二进制版本即可，即可进行普通的模型开发训练了，但如果要深入了解PyTorch根据官方文档，建议安装Python 3.7或以上的环境，而且需要C++14的编译器，比如clang，一开始我在ubuntu中装了Python的环境我也根据建议安装了Anaconda，一方面Anaconda会自动安装很多库，包括PyTorch所依赖的mk1这样的库。如果我们还需要编译支持GPU的PyTorch，需要安装cuda、cudnn，其中cuda建议安装10.2以上，cuDNN建议v7以上版本。另外，为了不影响本机环境，建议基于容器环境进行编译。

本机环境准备

笔者的开发环境是在一台比较老的PC机上，主机操作系统是Ubuntu18.04，配置了GPU卡GTX1660Ti。如果读者记不清

```
lspci | grep VGA
01:00.0 VGA compatible controller: NVIDIA Corporation Device 2182 (rev a1)
```

如果输出中没有GPU型号，如上面的输出，可以在以下网站查询得到：<http://pci-ids.ucw.cz/read/PC/10de/2182>

在确定GPU卡型号之后，可以在NVIDIA的网站上查找对应的驱动，网址为：<https://www.nvidia.com/Download/index.aspx?lang=en-us>。比如笔者的1660Ti的驱动信息如下：

```
> > Linux x64 (AMD64/EM64T) Display Driver >
> Version: 515.76 > Release Date: 2022.9.20 > Operating System: Linux
64-bit > Language: English (US) > File Size: 347.96 MB >
```

下载对应的驱动之后，安装即可。一般的电脑都有核心显卡，在安装的过程中可以考虑将核心显卡用于显示，独立显卡

如果是在主机环境编译，需要安装CUDA和Cudnn，根据NVIDIA官网的提示进行安装即可。

如果使用容器环境进行编译，本机还需要安装nvidia-container-runtime。

```
curl -s -L https://nvidia.github.io/nvidia-docker/gpgkey | sudo apt-key add -
distribution=$(. /etc/os-release;echo $ID$VERSION_ID)
echo $distribution
curl -s -L https://nvidia.github.io/nvidia-docker/$distribution/nvidia-docker.list | sudo tee /etc/apt/sources.list.d/nvidia-docker.list
#wget https://nvidia.github.io/nvidia-container-runtime/ubuntu14.04/amd64/.nvidia-container-runtime.list
sudo apt-get -y update
sudo apt-get install -y nvidia-container-toolkit
sudo apt-get install -y nvidia-container-runtime
sudo systemctl restart docker
```

之后需要安装docker，并将当前用户加入到docker的用户组里。

```
$ apt install docker.io
$ groupadd docker
$ usermod -ag docker <user>
```

在主机环境准备好后，我们开始准备基于ubuntu18.04的开放编译环境。

为了简便起见，建议直接使用NVIDIA预先准备好的容器环境，从这里可以找到对应本机操作系统和CUDA版本的容器：
<https://hub.docker.com/r/nvidia/cuda>。

比如笔者所使用的环境是Ubuntu18.04+CUDA11.7，因此应该使用的容器环境是：nvidia/cuda:11.7.0-cudnn8-devel-ubuntu18.04

启动容器的命令如下，读者朋友也可以根据需要加上其他的参数。笔者已经克隆了PyTorch的源码，放在\${HOME}/workspace/lab/

```
docker run -it --rm -v ${HOME}/workspace/lab:/lab --gpus all nvidia/cuda:11.7.0-cudnn8-devel-ubuntu18.04
```

另外，笔者编译PyTorch的时候，选择的是1.12.1的Tag，在编译的时候，要求cmake的版本高于3.13.0，而该容器自带的cmake版本是3.12.1，因此需要手动安装cmake。从官网下载cmake源代码，<https://cmake.org/download/>。解压后运行如下命令即可安装：

```
$ apt remove cmake
$ apt install libssl-dev
$ cd cmake-3.24.2
$ ./configure
$ make
$ make install
```

根据PyTorch README中的说明，需要在conda中安装多个依赖包：

```
$ conda install astunparse numpy ninja pyyaml setuptools cmake cffi typing_extensions future
$ conda install mkl mkl-include
```

编译步骤

```
$ git clone --recursive https://github.com/pytorch/pytorch
$ cd pytorch
# if you are updating an existing checkout
$ git submodule sync
$ git submodule update --init --recursive --jobs 0
$ git submodule update --init --recursive
```

启动容器，挂载PyTorch源码所在的目录，然后启动编译命令：

```
# DEBUG DEBUG=1 tools/setup_helpers/env.py -O0 -g'
# tools/setup_helpers/cmake.py make MAX_JOBS PC CPU
python setup.py clean
MAX_JOBS=2 DEBUG=1 USE_GPU=1 python setup.py build 2>&1 | tee build_test.log
```

在编译启动后，会创建build目录，之后所有的编译工作都在这个目录下完成。

如果没有什么问题，编译的最后输出如下：

```
-- Build files have been written to: /lab/tmp/pytorch/build
```

```
[1191/6244] Generating src/x86_64-fma/softmax.py.o
[1208/6244] Building C object confu-deps/XNNPACK/CMakeFiles/all_microkernels.dir/src/f32-dw
```

[1209/6244] Generating src/x86_64-fma/blas/sdotxf.py.o

.....

```
[ 0%] Linking C static library ../../../../lib/libclog.a
[ 0%] Linking C static library ../../lib/libpthreadpool.a
[ 1%] Linking CXX static library ../../../../lib/libgtestd.a
[ 2%] Linking C static library ../../../../lib/libtensorpipe_uv.a
[ 4%] Linking CXX static library ../../../../lib/libprotobuf-lited.a
[ 4%] Linking CXX static library ../../../../lib/libbenchmark.a
[ 4%] Linking CXX static library ../../../../lib/libgloo.a
[ 4%] Linking CXX static library ../../../../lib/libasmjit.a
[ 6%] Linking CXX static library ../../lib/libfmt.a
[ 7%] Linking CXX static library ../../../../lib/libprotobufd.a
[ 9%] Linking CXX shared library ../lib/libcaffe2_nvrtc.so
[ 9%] Linking CXX shared library ../lib/libc10.so
[ 9%] Linking C static library ../../lib/libfoxi_loader.a
[ 9%] Linking C executable ../../bin/mkrename
[ 9%] Linking C executable ../../bin/mkalias
[11%] Linking C executable ../../bin/mkdisp
[11%] Linking C shared library ../lib/libtorch_global_deps.so
[11%] Linking C executable ../../bin/mkrename_gnuabi
[11%] Linking C executable ../../bin/mkmasked_gnuabi
[11%] Linking C executable ../../bin/addSuffix
[13%] Linking C static library ../../lib/libcpuinfo.a
[15%] Linking C static library ../../lib/libcpuinfo_internals.a
[16%] Linking C static library ../../lib/libqnnpack.a
[16%] Linking C static library ../../lib/libnnpack_reference_layers.a
[18%] Linking CXX static library ../../lib/libpytorch_qnnpack.a
[23%] Linking CXX static library ../../../../lib/libprotocd.a
[23%] Linking CXX static library ../../../../lib/libbenchmark_main.a
[24%] Linking CXX static library ../../../../lib/libgtest_maind.a
[24%] Linking CXX static library ../../../../lib/libgmockd.a
[26%] Linking C static library ../../lib/libnnpack.a
[26%] Linking CXX static library ../../../../lib/libdnnl.a
[38%] Linking CXX static library ../../lib/libXNNPACK.a
[45%] Linking CXX static library ../../../../lib/libtensorpipe.a
[50%] Linking CXX executable ../../bin/c10_intrusive_ptr_benchmark
[51%] Linking CXX shared library ../lib/libc10_cuda.so
[54%] Linking CXX executable ../../bin/protoc
[54%] Linking CXX static library ../../../../lib/libkineto.a
[54%] Linking CXX static library ../../../../lib/libdnnl_graph.a
[54%] Linking CXX static library ../../../../lib/libgmock_maind.a
[56%] Linking C static library ../../lib/libsleep.a
[57%] Linking CXX static library ../../../../lib/libtensorpipe_cuda.a
[63%] Linking CXX static library ../../lib/libonnx_proto.a
```



```

[ 64%] Linking CXX static library ../lib/libcaffe2_protos.a
[ 70%] Linking CXX static library ../../lib/libonnx.a
[ 74%] Linking CXX static library ../../lib/libfbgemm.a
[ 74%] Linking CXX executable ../bin/vec_test_all_types_AVX2
[ 74%] Linking CXX executable ../bin/vec_test_all_types_DEFAULT
[ 88%] Linking CXX shared library ../lib/libtorch_cpu.so
Linking      libnccl.so.2.10.3                > /lab/pytorch-build/pytorch/build/nccl/lib/
[ 88%] Linking CXX static library ../../lib/libgloo_cuda.a
[ 93%] Linking CXX shared library ../lib/libtorch_cuda.so
[ 93%] Linking CXX shared library ../lib/libtorch.so
[ 93%] Linking CXX shared library ../lib/libtorch_cuda_linalg.so
[ 93%] Linking CXX executable ../bin/example_allreduce
[ 93%] Linking CXX executable ../bin/basic
[ 93%] Linking CXX executable ../bin/atest
[ 94%] Linking CXX executable ../bin/test_parallel
[ 94%] Linking CXX executable ../bin/verify_api_visibility
[ 94%] Linking CXX executable ../bin/mobile_memory_cleanup
[ 94%] Linking CXX shared library ../lib/libbackend_with_compiler.so
[ 94%] Linking CXX executable ../bin/tutorial_tensorexpr
[ 94%] Linking CXX shared library ../../../../lib/libshm.so
[ 94%] Linking CXX executable ../bin/parallel_benchmark
[ 95%] Linking CXX executable ../../../../bin/torch_shm_manager
[ 98%] Linking CXX executable ../bin/nvfuser_bench
[100%] Linking CXX shared library ../lib/libtorch_python.so
[100%] Linking CXX shared library ../lib/libnnapi_backend.so
building 'torch._C' extension

```

building 'torch._C_flatbuffer' extension

building 'torch._dl' extension

```

-----
|
|   It is no longer necessary to use the 'build' or 'rebuild' targets |
|
|   To install: |
|   $ python setup.py install |
|   To develop locally: |
|   $ python setup.py develop |
|   To force cmake to re-generate native build files (off by default): |
|   $ python setup.py develop --cmake |
|
|-----

```

PyTorch的setup.py

参考 https://blog.csdn.net/Sky_FULL1/article/details/125652654

PyTorch使用setuptools进行编译安装。

setuptools是常用的python库源码安装工具，其最主要的函数是setup(...)，所有安装包需要的参数包括包名。下面我们看一下PyTorch的setup.py，为了节约篇幅，并且考虑到绝大多数同学会使用Linux环境进行编译，这里删掉

```
# Constant known variables used throughout this file
cwd = os.path.dirname(os.path.abspath(__file__))
lib_path = os.path.join(cwd, "torch", "lib")
third_party_path = os.path.join(cwd, "third_party")
caffe2_build_dir = os.path.join(cwd, "build")

def configure_extension_build():
    #YL
    cmake_cache_vars = defaultdict(lambda: False, cmake.get_cmake_cache_variables())

    #YL

    library_dirs.append(lib_path)
    main_compile_args = []
    main_libraries = ['torch_python']
    main_link_args = []
    main_sources = ["torch/csrc/stub.c"]

    if cmake_cache_vars['USE_CUDA']:
        library_dirs.append(
            os.path.dirname(cmake_cache_vars['CUDA_CUDA_LIB']))

    if build_type.is_debug():
        extra_compile_args += ['-O0', '-g']
        extra_link_args += ['-O0', '-g']

#####
# Declare extensions and package
#####

extensions = []
packages = find_packages(exclude=('tools', 'tools.*'))
C = Extension("torch._C",
               libraries=main_libraries,
               sources=main_sources,
```

```

        language='c',
        extra_compile_args=main_compile_args + extra_compile_args,
        include_dirs=[],
        library_dirs=library_dirs,
        extra_link_args=extra_link_args + main_link_args + make_relative_rpath_args
C_flatbuffer = Extension("torch._C_flatbuffer",
                          libraries=main_libraries,
                          sources=["torch/csrc/stub_with_flatbuffer.c"],
                          language='c',
                          extra_compile_args=main_compile_args + extra_compile_args,
                          include_dirs=[],
                          library_dirs=library_dirs,
                          extra_link_args=extra_link_args + main_link_args + make_relative_rpath_args)
extensions.append(C)
extensions.append(C_flatbuffer)

if not IS_WINDOWS:
    DL = Extension("torch._dl",
                   sources=["torch/csrc/dl.c"],
                   language='c')
    extensions.append(DL)

# These extensions are built by cmake and copied manually in build_extensions()
# inside the build_ext implementation
if cmake_cache_vars['BUILD_CAFFE2']:
    extensions.append(
        Extension(
            name=str('caffe2.python.caffe2_pybind11_state'),
            sources=[]),
    )
    if cmake_cache_vars['USE_CUDA']:
        extensions.append(
            Extension(
                name=str('caffe2.python.caffe2_pybind11_state_gpu'),
                sources=[]),
        )
    if cmake_cache_vars['USE_ROCM']:
        extensions.append(
            Extension(
                name=str('caffe2.python.caffe2_pybind11_state_hip'),
                sources=[]),
        )

cmdclass = {
    'bdist_wheel': wheel_concatenate,
    'build_ext': build_ext,

```

```

        'clean': clean,
        'install': install,
        'sdist': sdist,
    }

    entry_points = ...

    return extensions, cmdclass, packages, entry_points, extra_install_requires

if __name__ == '__main__':
    extensions, cmdclass, packages, entry_points, extra_install_requires = configure_extensi
    setup(
        ext_modules=extensions,
        cmdclass=cmdclass,
        packages=packages,
        entry_points=entry_points,
        install_requires=install_requires,
        package_data={
            #YL
        },
        #YL
    )

```

PyTorch使用的是自定义的编译方法，指定了wheel_concatenate和build_ext这两个函数，分别负责库文件和扩展模

在编译库文件时，setuptools默认会编译打包以下文件：

- 由 py_modules 或 packages 指定的源文件
- 所有由 ext_modules 或 libraries 指定的 C 源码文件
- 由 scripts 指定的脚本文件
- 类似于 test/test*.py 的文件
- README.txt 或 README, setup.py, setup.cfg
- 所有 package_data 或 data_files 指定的文件

从上面的代码中可以看到，最主要的两个Extension是torch._C

基于cmake的编译体系

参考<https://blog.csdn.net/HaoBBNuanMM/article/details/115720457>

在build_ext()函数中，调用了Caffe2的编译，并且是在pytorch目录下开始编译的。

首先，打开开关CMAKE_EXPORT_COMPILE_COMMANDS，这样可以将所有的编译命令输出到一个文件里，我们可以在编译

```
set(CMAKE_EXPORT_COMPILE_COMMANDS ON)
```

之后设置优先使用CMake中的pthread库，据说libstdc++封装pthread库后，如果以dlopen的方式使用会导致空指针错
<https://zhuanlan.zhihu.com/p/128519905>

```
set(THREADS_PREFER_PTHREAD_FLAG ON)
```

之后是一些编译的配置，内容比较多，下面列出了一些主要的配置项。其中有很多配置项使用宏cmake_dependent_option(
cmake cmake_dependent_option(USE_CUDNN "Use cuDNN" ON "USE_CUDA"
OFF)代表当开启USE_CUDA的时候，也开启USE_CUDNN，否则关闭USE_CUDNN。

```
# ---[ Options.
# Note to developers: if you add an option below, make sure you also add it to
# cmake/Summary.cmake so that the summary prints out the option values.
include(CMakeDependentOption)
option(BUILD_BINARY "Build C++ binaries" OFF)
option(BUILD_PYTHON "Build Python binaries" ON)
option(BUILD_CAFFE2 "Master flag to build Caffe2" OFF)
cmake_dependent_option(
    BUILD_CAFFE2_OPS "Build Caffe2 operators" ON
    "BUILD_CAFFE2" OFF)
option(BUILD_SHARED_LIBS "Build libcaffe2.so" ON)
option(USE_CUDA "Use CUDA" ON)
cmake_dependent_option(
    USE_CUDNN "Use cuDNN" ON
    "USE_CUDA" OFF)
cmake_dependent_option(
    USE_NCCL "Use NCCL" ON
    "USE_CUDA OR USE_ROCM;UNIX;NOT APPLE" OFF)
option(USE_TENSORRT "Using Nvidia TensorRT library" OFF)

# Ensure that an MKLDNN build is the default for x86 CPUs
# but optional for AArch64 (dependent on -DUSE_MKLDNN).
cmake_dependent_option(
    USE_MKLDNN "Use MKLDNN. Only available on x86, x86_64, and AArch64." "${CPU_INTEL}"
    "CPU_INTEL OR CPU_AARCH64" OFF)

option(USE_DISTRIBUTED "Use distributed" ON)
cmake_dependent_option(
    USE_MPI "Use MPI for Caffe2. Only available if USE_DISTRIBUTED is on." ON
    "USE_DISTRIBUTED" OFF)
cmake_dependent_option(
    USE_GLOO "Use Gloo. Only available if USE_DISTRIBUTED is on." ON
    "USE_DISTRIBUTED" OFF)
```

PyTorch对ONNX的支持有两种方式，如果已有ONNX库，可以配置使用系统的自带的ONNX，否则重新编译生成。

```
if(NOT USE_SYSTEM_ONNX)
    set(ONNX_NAMESPACE "onnx_torch" CACHE STRING "A namespace for ONNX; needed to build with c
else()
    set(ONNX_NAMESPACE "onnx" CACHE STRING "A namespace for ONNX; needed to build with other f
endif()
```

接下来引用utils.cmake，这个文件里包括了很多工具函数，用于后边编译过程中的一些处理。

```
# ---[ Utils
include(cmake/public/Utils.cmake)
```

之后是一些版本号的设置，不再赘述。

这里设置了cmake的modules查找路径，以及编译输出的路径

```
# ---[ CMake scripts + modules
list(APPEND CMAKE_MODULE_PATH ${PROJECT_SOURCE_DIR}/cmake/Modules)

# ---[ CMake build directories
set(CMAKE_ARCHIVE_OUTPUT_DIRECTORY ${CMAKE_BINARY_DIR}/lib)
set(CMAKE_LIBRARY_OUTPUT_DIRECTORY ${CMAKE_BINARY_DIR}/lib)
set(CMAKE_RUNTIME_OUTPUT_DIRECTORY ${CMAKE_BINARY_DIR}/bin)
```

在编译的过程中，产生了下面这些动态库：

```
[ 2%] Linking C static library ../../lib/libtensorpipe_uv.a
[ 9%] Linking CXX shared library ../lib/libcaffe2_nvrtc.so
[ 9%] Linking CXX shared library ../lib/libc10.so
[ 11%] Linking C shared library ../lib/libtorch_global_deps.so
[ 45%] Linking CXX static library ../../lib/libtensorpipe.a
[ 51%] Linking CXX shared library ../lib/libc10_cuda.so
[ 57%] Linking CXX static library ../../lib/libtensorpipe_cuda.a
[ 88%] Linking CXX shared library ../lib/libtorch_cpu.so
Linking      libnccl.so.2.10.3                > /lab/pytorch-build/pytorch/build/nccl/lib/
[ 93%] Linking CXX shared library ../lib/libtorch_cuda.so
[ 93%] Linking CXX shared library ../lib/libtorch.so
[ 93%] Linking CXX shared library ../lib/libc10d_cuda_test.so
[ 93%] Linking CXX shared library ../lib/libtorch_cuda_linalg.so
[ 93%] Linking CXX executable ../bin/NamedTensor_test
[ 94%] Linking CXX executable ../bin/scalar_tensor_test
[ 94%] Linking CXX executable ../bin/undefined_tensor_test
[ 94%] Linking CXX executable ../bin/lazy_tensor_test
[ 94%] Linking CXX executable ../bin/tensor_iterator_test
[ 94%] Linking CXX executable ../bin/cuda_packedtensoraccessor_test
[ 94%] Linking CXX shared library ../lib/libjitbackend_test.so
[ 94%] Linking CXX shared library ../lib/libtorchbind_test.so
[ 94%] Linking CXX shared library ../lib/libbackend_with_compiler.so
[ 94%] Linking CXX executable ../bin/tutorial_tensorexpr
[ 94%] Linking CXX shared library ../../lib/libshm.so
[ 98%] Linking CXX executable ../bin/test_tensorexpr
[100%] Linking CXX shared library ../lib/libtorch_python.so
[100%] Linking CXX shared library ../lib/libnnapi_backend.so
```

最后，在通过cmake将必要的库编译完成以后，再执行setup.py中的编译命令，生成PyTorch所依赖的扩展：

```
building 'torch._C' extension
creating build/temp.linux-x86_64-3.9
```

```

creating build/temp.linux-x86_64-3.9/torch
creating build/temp.linux-x86_64-3.9/torch/csrc
gcc -pthread -B /root/anaconda3/compiler_compat -Wno-unused-result -Wsign-compare -DNDEBUG -
gcc -pthread -B /root/anaconda3/compiler_compat -shared -Wl,-rpath,/root/anaconda3/lib -Wl,-
building 'torch._C_flatbuffer' extension
gcc -pthread -B /root/anaconda3/compiler_compat -Wno-unused-result -Wsign-compare -DNDEBUG -
gcc -pthread -B /root/anaconda3/compiler_compat -shared -Wl,-rpath,/root/anaconda3/lib -Wl,-
building 'torch._dl' extension
gcc -pthread -B /root/anaconda3/compiler_compat -Wno-unused-result -Wsign-compare -DNDEBUG -
gcc -pthread -B /root/anaconda3/compiler_compat -shared -Wl,-rpath,/root/anaconda3/lib -Wl,-

```

对比着，在安装pytorch后，我们可以看到torch目录下有如下的动态库：

```

./_dl.cpython-36m-x86_64-linux-gnu.so
./lib/libtorch_python.so
./lib/libcaffe2_observers.so
./lib/libcaffe2_nvrtc.so
./lib/libc10.so
./lib/libc10_cuda.so
./lib/libshm.so
./lib/libcaffe2_detectron_ops_gpu.so
./lib/libtorch.so
./lib/libcaffe2_module_test_dynamic.so
./_C.cpython-36m-x86_64-linux-gnu.so

```

Caffe2下有列动态库：“ ‘Bash ./python/caffe2_pybind11_state.cpython-36m-x86_64-linux-gnu.so ./python/caffe2_pybind11_state_gpu.cpython-36m-x86_64-linux-gnu.so . . .

```

# ---[ Misc checks to cope with various compiler modes
include(cmake/MiscCheck.cmake)

# External projects
include(ExternalProject)

include(cmake/Dependencies.cmake)

# ---[ Allowlist file if allowlist is specified
include(cmake/Allowlist.cmake)

# Prefix path to Caffe2 headers.
# If a directory containing installed Caffe2 headers was inadvertently
# added to the list of include directories, prefixing
# PROJECT_SOURCE_DIR means this source tree always takes precedence.
include_directories(BEFORE ${PROJECT_SOURCE_DIR})

```

```

# Prefix path to generated Caffe2 headers.
# These need to take precedence over their empty counterparts located
# in PROJECT_SOURCE_DIR.
include_directories(BEFORE ${PROJECT_BINARY_DIR})

include_directories(BEFORE ${PROJECT_SOURCE_DIR}/aten/src/)
include_directories(BEFORE ${PROJECT_BINARY_DIR}/aten/src/)

# ---[ Main build
add_subdirectory(c10)
add_subdirectory(caffe2)

# ---[ Modules
# If master flag for building Caffe2 is disabled, we also disable the
# build for Caffe2 related operator modules.
if(BUILD_CAFFE2)
    add_subdirectory(modules)
endif()

# ---[ Binaries
# Binaries will be built after the Caffe2 main libraries and the modules
# are built. For the binaries, they will be linked to the Caffe2 main
# libraries, as well as all the modules that are built with Caffe2 (the ones
# built in the previous Modules section above).
if(BUILD_BINARY)
    add_subdirectory(binaries)
endif()

include(cmake/Summary.cmake)
caffe2_print_configuration_summary()

# ---[ Torch Deploy
if(USE_DEPLOY)
    add_subdirectory(torch/csrc/deploy)
endif()

```

PyTorch 动态代码生成

参考 <https://zhuanlan.zhihu.com/p/59425970> 参考 <https://zhuanlan.zhihu.com/p/55966063>

PyTorch代码主要包括三部分：

- C10. C10是Caffe Tensor Library的缩写。PyTorch目前正在将代码从ATen/core目录迁移到C10。
- ATen, ATen是A TENSOR library for C++11的缩写，是PyTorch的C++ tensor library。ATen部分有大量的代码是来声明和定义Tensor运算相关的逻辑的，除此之外，PyTorch还使用了aten/src/ATen目录下的代码。
- Torch, 部分代码仍然在使用以前的快要进入历史博物馆的Torch开源项目，比如具有下面这些文件名格式的文件：


```

TH* = Torch
THC* = Torch Cuda
THCS* = Torch Cuda Sparse (now defunct)
THCUNN* = Torch CUda Neural Network (see cunn)
THD* = Torch Distributed
THNN* = Torch Neural Network
THS* = Torch Sparse (now defunct)
THP* = Torch Python

```

PyTorch会使用tools/setup_helpers/generate_code.py来动态生成Torch层面相关的一些代码，这部分动态生成的代码C10目前最具代表性的一个class就是TensorImpl了，它实现了Tensor的最基础框架。继承者和使用者有：

编译第三方的库

```

#Facebook  cpuinfo  cpu
third_party/cpuinfo

#Facebook
# Pytorch caffe2 ncnn coreml
third_party/onnx

#FB (Facebook) + GEMM (General Matrix-Matrix Multiplication)
#Facebook      caffe2 x86      backend
third_party/fbgemm

#  benchmark
third_party/benchmark

#  protobuf
third_party/protobuf

#  UT
third_party/googletest

#Facebook
third_party/QNNPACK

#
third_party/gloo

#Intel  MKL-DNN
third_party/ideep

```

代码生成

ATen的native函数是PyTorch目前主推的operator机制，作为对比，老旧的TH/THC函数（使用cwrap定义）将逐渐被Aop需要修改这个yaml文件。

“ ‘

代码生成相关的工具在tools目录下：

```
autograd
  gen_annotated_fn_args.py
  gen_autograd_functions.py
  gen_autograd.py
  gen_inplace_or_view_type.py
  gen_python_functions.py
  gen_trace_type.py
  gen_variable_factories.py
  gen_variable_type.py
  templates
    ADInplaceOrViewType.cpp
    annotated_fn_args.py.in
    Functions.cpp
    Functions.h
    python_fft_functions.cpp
    python_functions.cpp
    python_functions.h
    python_linalg_functions.cpp
    python_nn_functions.cpp
    python_return_types.cpp
    python_sparse_functions.cpp
    python_special_functions.cpp
    python_torch_functions.cpp
    python_variable_methods.cpp
    TraceType.cpp
    variable_factories.h
    VariableType.cpp
    VariableType.h
code_analyzer
  gen_operators_yaml.py
  gen_oplist.py
  gen_op_registration_allowlist.py
generated_dirs.txt
jit
  gen_unboxing.py
  templates
    aten_schema_declarations.cpp
    external_functions_codegen_template.cpp
```

```

setup_helpers
  generate_code.py
  gen.py
  gen_unboxing.py
  gen_version_header.py

```

我们先看几个重要的文件：

- generated_dirs.txt: 这个文件里列举了编译过程中自动生成的代码所在的路径，当前版本中该文件的内容如下：


```

torch/csrc/autograd/generated/      #
torch/csrc/jit/generated/           # JIT
build/aten/src/ATen                  # aten

```
- setup_helpers/generate_code.py: 这个文件中函数generate_code()是代码生成的入口。等下我们会沿着这个函数...

代码生成的流程

代码生成沿着以下的流程进行：

调用tools/autograd/gen_autograd.py中的函数gen_autograd_python，这个函数输入参数
 NATIVE_FUNCTIONS_PATH = “aten/src/ATen/native/native_functions.yaml”
 TAGS_PATH = “aten/src/ATen/native/tags.yaml”

```

<li> native_functions_path: native functions
<li> derivatives.yaml:
<li> templates:
<li> deprecated.yaml:
</ol>

gen_python_functions.gen()      ATen Python torch._C nn_fft_linalg_sparse_special
<ol>
<li> native_functions.yaml tags.yaml native
<li>
<li> deprecated.yaml
<li> FileManager.write_with_template() python_variable_methods.cpp
    python_variable_methods

#define TORCH_ASSERT_ONLY_METHOD_OPERATORS
// ${generated_comment}

// ...
#include <stdexcept>

#ifdef AT_PER_OPERATOR_HEADERS
#include <ATen/Functions.h>
#else
$ops_headers

```

```

#endif

//...

// generated methods start here

${py_methods}

static PyObject * THPVariable_bool_scalar(PyObject* self, PyObject* args) {

//...

    {"tolist", THPVariable_tolist, METH_NOARGS, NULL},
    {"type", castPyCFunctionWithKeywords(THPVariable_type), METH_VARARGS | METH_KEYWORDS, NULL},
    ${py_method_defs}
    {NULL}
};

```

其中的关键变量是py_methods，这个变量包含了很多函数的定义，其中每个函数是根据模板字符串生成的，如下是其

// tools/autograd/gen_python_functions.py

```

PY_VARIABLE_METHOD_VARARGS = CodeTemplate(
    r"""
// ${name}
static PyObject * ${pycname}(PyObject* self_, PyObject* args, PyObject* kwargs)
{
    ${method_header}
    static PythonArgParser parser({
        ${signatures}
    }, /*traceable=*/${traceable});

    ParsedArgs<${max_args}> parsed_args;
    auto _r = parser.parse(${self_}, args, kwargs, parsed_args);
    ${check_has_torch_function}
    switch (_r.idx) {
        ${dispatch}
    }
    ${method_footer}
}

"""
)

```

根据这个模板生成的函数代码大概是下面这样：

// torch/csrc/autograd/generated/python_variable_methods.cpp

```

static PyObject * THPVariable_add(PyObject* self_, PyObject* args, PyObject* kwargs)
{
    HANDLE_TH_ERRORS
    const Tensor& self = THPVariable_Unpack(self_);
    static PythonArgParser parser({
        "add(Scalar alpha, Tensor other)|deprecated",
        "add(Tensor other, *, Scalar alpha=1)",
    }, /*traceable=*/true);

    ParsedArgs<2> parsed_args;
    auto _r = parser.parse(self_, args, kwargs, parsed_args);
    if(_r.has_torch_function()) {
        return handle_torch_function(_r, self_, args, kwargs, THPVariableClass, "torch.Tensor");
    }
    switch (_r.idx) {
    case 0: {
        // [deprecated] aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor

        auto dispatch_add = [](const at::Tensor & self, const at::Scalar & alpha, const at::Tensor & other) {
            pybind11::gil_scoped_release no_gil;
            return self.add(other, alpha);
        };
        return wrap(dispatch_add(self, _r.scalar(0), _r.tensor(1)));
    }
    case 1: {
        // aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor

        auto dispatch_add = [](const at::Tensor & self, const at::Tensor & other, const at::Scalar & alpha) {
            pybind11::gil_scoped_release no_gil;
            return self.add(other, alpha);
        };
        return wrap(dispatch_add(self, _r.tensor(0), _r.scalar(1)));
    }
    }
    Py_RETURN_NONE;
    END_HANDLE_TH_ERRORS
}

```


生成的库

```

# /pytorch/build/lib.linux-x86_64-3.7/torch
./_C.cpython-37m-x86_64-linux-gnu.so

```

```

./lib/libtorch_python.so
./lib/libtorchbind_test.so
./lib/libtorch_cpu.so
./lib/libjitbackend_test.so
./lib/libc10.so
./lib/libshm.so
./lib/libtorch.so
./lib/libtorch_global_deps.so
./lib/libbackend_with_compiler.so
./_C_flatbuffer.cpython-37m-x86_64-linux-gnu.so
./_dl.cpython-37m-x86_64-linux-gnu.so

```

其中_C.cpython-37m-x86_64-linux-gnu.so是主要的入口点，后面的章节我们会从这个入口点分析PyTorch的初始化（found可忽略）。

```
# pytorch/build/lib.linux-x86_64-3.7/torch
```

```

$ ldd ./_C.cpython-37m-x86_64-linux-gnu.so
    linux-vdso.so.1 (0x0000ffff18175000)
    libtorch_python.so => /home/harry/lab/tmp/pytorch/build/lib.linux-x86_64-3.7/torch/./lib/libtorch_python.so
    libpthread.so.0 => /lib/x86_64-linux-gnu/libpthread.so.0 (0x00007feff2b42000)
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007feff2751000)
    libshm.so => /home/harry/lab/tmp/pytorch/build/lib.linux-x86_64-3.7/torch/./lib/libshm.so
    libtorch.so => /home/harry/lab/tmp/pytorch/build/lib.linux-x86_64-3.7/torch/./lib/libtorch.so
    libtorch_cpu.so => /home/harry/lab/tmp/pytorch/build/lib.linux-x86_64-3.7/torch/./lib/libtorch_cpu.so
    libc10.so => /home/harry/lab/tmp/pytorch/build/lib.linux-x86_64-3.7/torch/./lib/libc10.so
    libstdc++.so.6 => /usr/lib/x86_64-linux-gnu/libstdc++.so.6 (0x00007fefddc7c000)
    libm.so.6 => /lib/x86_64-linux-gnu/libm.so.6 (0x00007fefdd8de000)
    libgcc_s.so.1 => /lib/x86_64-linux-gnu/libgcc_s.so.1 (0x00007fefdd6c6000)
    /lib64/ld-linux-x86-64.so.2 (0x00007feff4fcc000)
    librt.so.1 => /lib/x86_64-linux-gnu/librt.so.1 (0x00007fefdd4be000)
    libgomp.so.1 => /usr/lib/x86_64-linux-gnu/libgomp.so.1 (0x00007fefdd28f000)
    libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007fefdd08b000)
    libmkl_intel_lp64.so => not found
    libmkl_gnu_thread.so => not found
    libmkl_core.so => not found

```

常见问题

- submodule没有下载完整 一个简单的处理办法是删除third_party下的相关目录，然后手动git clone即可。相关的git url定义在.submodule以及.gi/config中
- 编译时出现RPATH相关的问题 处理办法是先运行clean命令，然后再编译

```

> python setup.py clean
> python setup.py build

```

- lib库找不到 错误详情: No rule to make target ' /usr/lib/x86_64-linux-gnu/libXXX.so ' ‘bash > find / -name “librt.so.*” > ln -s /lib/x86_64-linux-gnu/librt.so.1 /usr/lib/x86_64-linux-gnu/librt.so
- c++ bash > apt install g++ “ ‘ 注意, 如果安装clang, 也可以编译, 但c++的版本如果比较低, 比如6.0, 就命令编译开关没找到 的问题。

- 在PC上编译时Hang住

一般来说为了加快编译速度, 编译大型项目时都会采用并行编译的方式, pytorch的编译也是, 启动编译后, 可以在简单起见, 在启动编译前, 可以设置环境变量CMAKE_BUILD_PARALLEL_LEVEL来减少编译的并行度。

- 编译Debug版本时出现internal compiler error

如果只是在编译Debug版本时出现, 可能是和优化编译选项有冲突, 因为优化编译选项-O1 -O2 -O3可能会重新排列代码导致代码对应出现问题, 排查真正的问题非常困难, 建议简单处理, 对出现问题的编译选项或者-O 选项。

PyTorch的编译由setup.py发起, 但真正执行编译时, 相关的命令写在build/build.ninja里, 只要在这个文件里修改

参考

<https://zhuanlan.zhihu.com/p/321449610>

<https://blog.5lcto.com/SpaceVision/5072093>

<https://zhuanlan.zhihu.com/p/55204134>

<https://github.com/pytorch/pytorch#from-source>

从零开始编译PyTorch软件包 <https://zhuanlan.zhihu.com/p/347084475>

Pytorch setup.py 详解 https://blog.csdn.net/Sky_FULL1/article/details/125652654

PyTorch 动态代码生成 <https://zhuanlan.zhihu.com/p/55966063>

PyTorch 动态代码生成 <https://zhuanlan.zhihu.com/p/59425970>

<https://blog.csdn.net/HaoBBNuanMM/article/details/115720457>

PyTorch引擎的主要模块及初始化

主要内容

本章对PyTorch的整体架构做了初步的分析, 这部分也是理解PyTorch核心引擎工作机制的关键部分, 在这里我们力图PyTorch从上层到C++的底层包括哪些重要的模块

这些模块是如何初始化的

从设计上看, 这些模块是如何配合的

PyTorch的核心模块

- PythonAPI
- C++部分Engine
- THP
- ATen
- JITWdq

```
src
!--- ATen      # Tensor  C++
|--- TH        # Tensor CPU
|--- THC       # Tensor CUDA
|--- THCUNN    #   CUDA
|--- THNN      #   CPU

torch
|--- csrc      # Torch C++
    |--- module.cpp      # Torch C++
```

PyTorch的C++扩展模块初始化

C++扩展模块_C可以说是PyTorch的核心，是PyTorch代码量最大最复杂的部分，下面我们来看看这个模块是如何加载

C++扩展模块的加载

在加载torch模块的时候，python会执行torch/init.py. 其中会加载_C模块，根据Python3的规范，如果某个模块是.so，在linux环境下，对应的就是_C.cpython-37m-x86_64-linux-gnu.so。

加载这个动态库后，会调用其中的initModule()函数。在这个函数中，进行了一系列的初始化工作

```
// torch/csrc/Module.cpp
```

```
PyObject* initModule() {

    // ...

    // _C
    THPUutils_addPyMethodDefs(methods, TorchMethods);
    THPUutils_addPyMethodDefs(methods, DataLoaderMethods);
    THPUutils_addPyMethodDefs(methods, torch::autograd::python_functions());
    THPUutils_addPyMethodDefs(methods, torch::multiprocessing::python_functions());

    THPUutils_addPyMethodDefs(methods, THCPModule_methods());
```



```

THPUUtils_addPyMethodDefs(methods, torch::distributed::c10d::python_functions());

THPUUtils_addPyMethodDefs(methods, torch::distributed::rpc::python_functions());
THPUUtils_addPyMethodDefs(
    methods, torch::distributed::autograd::python_functions());
THPUUtils_addPyMethodDefs(methods, torch::distributed::rpc::testing::python_functions());

//  _C
static struct PyModuleDef torchmodule = {
    PyModuleDef_HEAD_INIT,
    "torch._C",
    nullptr,
    -1,
    methods.data()
};
ASSERT_TRUE(module = PyModule_Create(&torchmodule));
ASSERT_TRUE(THPGenerator_init(module));
ASSERT_TRUE(THPException_init(module));
THPSize_init(module);
THPDtype_init(module);
THPDTypeInfo_init(module);
THPLayout_init(module);
THPMemoryFormat_init(module);
THPQScheme_init(module);
THPDevice_init(module);
THPStream_init(module);

//  Tensor
ASSERT_TRUE(THPVariable_initModule(module));
ASSERT_TRUE(THPFunction_initModule(module));
ASSERT_TRUE(THPEngine_initModule(module));
// NOTE: We need to be able to access OperatorExportTypes from ONNX for use in
// the export side of JIT, so this ONNX init needs to appear before the JIT
// init.
torch::onnx::initONNXBindings(module);
torch::jit::initJITBindings(module);
torch::monitor::initMonitorBindings(module);
torch::impl::dispatch::initDispatchBindings(module);
torch::throughput_benchmark::initThroughputBenchmarkBindings(module);
torch::autograd::initReturnTypes(module);
torch::autograd::initNNFunctions(module);
torch::autograd::initFFTFunctions(module);
torch::autograd::initLinalgFunctions(module);
torch::autograd::initSparseFunctions(module);
torch::autograd::initSpecialFunctions(module);
torch::autograd::init_legacy_variable(module);

```

```

torch::python::init_bindings(module);
torch::lazy::initLazyBindings(module);
#ifdef USE_CUDA
torch::cuda::initModule(module);
#endif
ASSERT_TRUE(THPStorage_init(module));

#ifdef USE_CUDA
// This will only initialise base classes and attach them to library namespace
// They won't be ready for real usage until importing cuda module, that will
// complete the process (but it defines Python classes before calling back into
// C, so these lines have to execute first)..
THCPStream_init(module);
THCPEvent_init(module);
THCPGraph_init(module);
#endif

auto set_module_attr = [&](const char* name, PyObject* v, bool incref = true) {
    // PyModule_AddObject steals reference
    if (incref) {
        Py_INCREF(v);
    }
    return PyModule_AddObject(module, name, v) == 0;
};

// ...

ASSERT_TRUE(set_module_attr("has_openmp", at::hasOpenMP() ? Py_True : Py_False));
ASSERT_TRUE(set_module_attr("has_mkl", at::hasMKL() ? Py_True : Py_False));
ASSERT_TRUE(set_module_attr("has_lapack", at::hasLAPACK() ? Py_True : Py_False));

// ...

py::enum_<at::native::ConvBackend>(py_module, "_ConvBackend")
    .value("CudaDepthwise2d", at::native::ConvBackend::CudaDepthwise2d)
    .value("CudaDepthwise3d", at::native::ConvBackend::CudaDepthwise3d)
    .value("Cudnn", at::native::ConvBackend::Cudnn)
    .value("CudnnTranspose", at::native::ConvBackend::CudnnTranspose)
    .value("Empty", at::native::ConvBackend::Empty)
    .value("Miopen", at::native::ConvBackend::Miopen)
    .value("MiopenDepthwise", at::native::ConvBackend::MiopenDepthwise)
    .value("MiopenTranspose", at::native::ConvBackend::MiopenTranspose)
    .value("Mkldnn", at::native::ConvBackend::Mkldnn)
    .value("MkldnnEmpty", at::native::ConvBackend::MkldnnEmpty)
    .value("NnpackSpatial", at::native::ConvBackend::NnpackSpatial)
    .value("Overrideable", at::native::ConvBackend::Overrideable)

```

```

        .value("Slow2d", at::native::ConvBackend::Slow2d)
        .value("Slow3d", at::native::ConvBackend::Slow3d)
        .value("SlowDilated2d", at::native::ConvBackend::SlowDilated2d)
        .value("SlowDilated3d", at::native::ConvBackend::SlowDilated3d)
        .value("SlowTranspose2d", at::native::ConvBackend::SlowTranspose2d)
        .value("SlowTranspose3d", at::native::ConvBackend::SlowTranspose3d)
        .value("Winograd3x3Depthwise", at::native::ConvBackend::Winograd3x3Depthwise)
        .value("Xnnpack2d", at::native::ConvBackend::Xnnpack2d);

py_module.def("_select_conv_backend", [] (
    const at::Tensor& input, const at::Tensor& weight, const c10::optional<at::Tensor>&
    at::IntArrayRef stride_, at::IntArrayRef padding_, at::IntArrayRef dilation_,
    bool transposed_, at::IntArrayRef output_padding_, int64_t groups_) {
    return at::native::select_conv_backend(
        input, weight, bias_opt, stride_, padding_, dilation_, transposed_, output_padding_);
});

py::enum_<at::LinalgBackend>(py_module, "_LinalgBackend")
    .value("Default", at::LinalgBackend::Default)
    .value("Cusolver", at::LinalgBackend::Cusolver)
    .value("Magma", at::LinalgBackend::Magma);

py_module.def("_set_linalg_preferred_backend", [] (at::LinalgBackend b) {
    at::globalContext().setLinalgPreferredBackend(b);
});
py_module.def("_get_linalg_preferred_backend", [] () {
    return at::globalContext().linalgPreferredBackend();
});

// ...

return module;
END_HANDLE_TH_ERRORS
}

```

Torch 函数库的初始化

在Python层面，模块torch提供了非常多的函数，比如torch.abs()，torch.randn()，torch.ones()等等，在初始化_C模块的时候，这些函数也被注册到_C模块中。

```

// torch/csrc/autograd/python_variable.cpp

bool THPVariable_initModule(PyObject *module)
{
    // ...
}

```

```

PyModule_AddObject(module, "_TensorBase", (PyObject *)&THPVariableType);
torch::autograd::initTorchFunctions(module);
// ...
return true;
}

```

在下面的代码中，我们可以看到，相关的函数被收集到torch_functions中，同时这个函数列表也被注册到_C的_Var

```

// torch/csrc/autograd/python_torch_functions_manual.cpp

void initTorchFunctions(PyObject *module) {
    static std::vector<PyMethodDef> torch_functions;
    gatherTorchFunctions(torch_functions);
    THPVariableFunctions.tp_methods = torch_functions.data();

    //...
    if (PyModule_AddObject(module, "_VariableFunctionsClass",
                           reinterpret_cast<PyObject*>(&THPVariableFunctions)) < 0) {
        throw python_error();
    }
    // PyType_GenericNew returns a new reference
    THPVariableFunctionsModule = PyType_GenericNew(&THPVariableFunctions, Py_None, Py_None);
    // PyModule_AddObject steals a reference
    if (PyModule_AddObject(module, "_VariableFunctions", THPVariableFunctionsModule) < 0) {
        throw python_error();
    }
}

```

在torch模块的初始化过程中，_C模块的子模块_VariableFunctions中的所有属性都被注册到torch模块中，当然也有

```

# torch/__init__.py

for name in dir(_C._VariableFunctions):
    if name.startswith('__') or name in PRIVATE_OPS:
        continue
    obj = getattr(_C._VariableFunctions, name)
    obj.__module__ = 'torch'
    globals()[name] = obj
    if not name.startswith("_"):
        __all__.append(name)

```

下面我们看看具体有哪些函数被注册了。函数列表是通过gatherTorchFunctions()来收集的，这个函数又调用了gatherTorchFunctions_1(), gatherTorchFunctions_2()这几个函数。

```

// torch/csrc/autograd/python_torch_functions_manual.cpp

void gatherTorchFunctions(std::vector<PyMethodDef> &torch_functions) {
    constexpr size_t num_functions = sizeof(torch_functions_manual) / sizeof(torch_functions_m
    torch_functions.assign(torch_functions_manual,

```

```

        torch_functions_manual + num_functions);
// NOTE: Must be synced with num_shards in tools/autograd/gen_python_functions.py
gatherTorchFunctions_0(torch_functions);
gatherTorchFunctions_1(torch_functions);
gatherTorchFunctions_2(torch_functions);

//...

```

为什么这样设计呢？大概有两个原因：- 函数的数量很多，而且在不断的增加，需要方便扩展
- 函数大多是算子，算子和平台相关，每个算子有多种实现，同样为了在不同平台迁移扩展，PyTorch设计了代码生成器gatherTorchFunctions_N()这几个函数是通过模板生成的，完成编译后，可以在下面的文件中找到：

```

// torch/csrc/autograd/generated/python_torch_functions_0.cpp

static PyMethodDef torch_functions_shard[] = {
    {"_cast_Byte", castPyCFunctionWithKeywords(THPVariable__cast_Byte), METH_VARARGS | METH_KEYWORDS},
    //...
    {"eye", castPyCFunctionWithKeywords(THPVariable_eye), METH_VARARGS | METH_KEYWORDS | METH_NOARGS},
    {"rand", castPyCFunctionWithKeywords(THPVariable_rand), METH_VARARGS | METH_KEYWORDS | METH_NOARGS},
    //...
};

void gatherTorchFunctions_0(std::vector<PyMethodDef> &torch_functions) {
    constexpr size_t num_functions = sizeof(torch_functions_shard) / sizeof(torch_functions_shard[0]);
    torch_functions.insert(
        torch_functions.end(),
        torch_functions_shard,
        torch_functions_shard + num_functions);
}

```

Tensor

在Pytorch的早期版本中，Tensor被定义在TH模块中的THTensor类中，后来TH模块被移除了，也就有了更直观的Tensor。当前Tensor的定义在TensorBody.h中，

```

// torch/include/ATen/core/TensorBody.h

class TORCH_API Tensor: public TensorBase {
public:
    Tensor(const Tensor &tensor) = default;
    Tensor(Tensor &&tensor) = default;

    using TensorBase::size;
    using TensorBase::stride;
}

```

```

Tensor cpu() const {
    return to(options().device(DeviceType::CPU), /*non_blocking*/ false, /*copy*/ false);
}

// TODO: The Python version also accepts arguments
Tensor cuda() const {
    return to(options().device(DeviceType::CUDA), /*non_blocking*/ false, /*copy*/ false);
}

void backward(const Tensor & gradient={}, ...) const {
    ...
}
}

```

我们还可以看到，Tensor类本身的实现很少，大部分功能来自于其父类TensorBase。根据文档注释我们可以了解到，

// torch/include/ATen/core/TensorBase.h

```

class TORCH_API TensorBase {

    int64_t dim() const {
        return impl_->dim();
    }
    int64_t storage_offset() const {
        return impl_->storage_offset();
    }

    // ...

    bool requires_grad() const {
        return impl_->requires_grad();
    }
    bool is_leaf() const;
    TensorBase data() const;

    c10::intrusive_ptr<TensorImpl, UndefinedTensorImpl> impl_;
}

```

https://blog.csdn.net/Chris_zhangrx/article/details/119086815

c10::intrusive_ptr是PyTorch的内部智能指针实现，其工作方式如下：

首先完美转发所有的参数来构建 intrusive_ptr 用这些参数 new 一个新的 TTarget 类型的对象 用新的 TTarget 对象构造一个 intrusive_ptr 构造 intrusive_ptr 的同时对 refcount_ 和 weakcount_ 都加 1，如果是默认构造，则两个引用计数都默认为 0，根据这个可以将通过 make_intrusive 构造的指针与堆栈上的会被自动析构的情况分开，用来确保内存是我们自己分配的。

以后有机会我们再研究一下intrusive_ptr的实现，在此之前，我们主要关注impl_这个成员变量，也就是TensorImpl

```
// c10/core/TensorImpl.h
```

```
struct C10_API TensorImpl : public c10::intrusive_ptr_target {

TensorImpl(
    Storage&& storage,
    DispatchKeySet,
    const caffe2::TypeMeta data_type);

public:
    TensorImpl(const TensorImpl&) = delete;
    TensorImpl& operator=(const TensorImpl&) = delete;
    TensorImpl(TensorImpl&&) = delete;
    TensorImpl& operator=(TensorImpl&&) = delete;

    DispatchKeySet key_set() const {
        return key_set_;
    }

    int64_t dim() const {
        //...
    }
    bool is_contiguous(
        //...
    )

    Storage storage_;

private:
    std::unique_ptr<c10::AutogradMetaInterface> autograd_meta_ = nullptr;

protected:
    std::unique_ptr<c10::NamedTensorMetaInterface> named_tensor_meta_ = nullptr;

    c10::VariableVersion version_counter_;

    std::atomic<impl::PyInterpreter*> pyobj_interpreter_;
    PyObject* pyobj_;

    c10::impl::SizesAndStrides sizes_and_strides_;

    int64_t storage_offset_ = 0;
    int64_t numel_ = 1;
```

```

caffe2::TypeMeta data_type_;
c10::optional<c10::Device> device_opt_;

bool is_contiguous_ : 1;

bool storage_access_should_throw_ : 1;

bool is_channels_last_ : 1;
bool is_channels_last_contiguous_ : 1;
bool is_channels_last_3d_ : 1;

bool is_channels_last_3d_contiguous_ : 1;

bool is_non_overlapping_and_dense_ : 1;

bool is_wrapped_number_ : 1;

bool allow_tensor_metadata_change_ : 1;

bool reserved_ : 1;
uint8_t sizes_strides_policy_ : 2;

DispatchKeySet key_set_;
}

```

对于TensorImpl类来说，比较重要的成员变量有以下几个： - storage_。这个变量存储了真正的张量数据 - autograd_meta_。存储反向传播所需要的元信息，如梯度计算函数和梯度等。 - pyobj_。Tensor所对应的Python Object - data_type_。Tensor内的数据类型。 - device_opt_。存放Tensor的设备。 -

下面我们看一下Tensor的存储，因为Tensor的存储方式和算子的计算息息相关，对性能的影响也非常的关键。

// c10/core/Storage.h

```

struct C10_API Storage {
    //...

protected:
    c10::intrusive_ptr<StorageImpl> storage_impl_;
}

```

和Tensor的定义类似，Storage也是使用StorageImpl类来隐藏其复杂的实现。因此我们主要关注StorageImpl的实现

// c10/core/StorageImpl.h

```

struct C10_API StorageImpl : public c10::intrusive_ptr_target {
public:

```



```

struct use_byte_size_t {};

StorageImpl(
    use_byte_size_t /*use_byte_size*/,
    size_t size_bytes,
    at::DataPtr data_ptr,
    at::Allocator* allocator,
    bool resizable)
    : data_ptr_(std::move(data_ptr)),
      size_bytes_(size_bytes),
      resizable_(resizable),
      received_cuda_(false),
      allocator_(allocator) {
    if (resizable) {
        TORCH_INTERNAL_ASSERT(
            allocator_, "For resizable storage, allocator must be provided");
    }
}

void* data() {
    return data_ptr_.get();
}

at::DeviceType device_type() const {
    return data_ptr_.device().type();
}

private:
    DataPtr data_ptr_;
    size_t size_bytes_;
    bool resizable_;
    // Identifies that Storage was received from another process and doesn't have
    // local to process cuda memory allocation
    bool received_cuda_;
    Allocator* allocator_;
}

StorageImpl的关键成员是data_ptr_，其定义在这里：
// c10/core/Allocator.h

class C10_API DataPtr {
private:
    c10::detail::UniqueVoidPtr ptr_;
    Device device_;
}

```

```
// c10/util/UniqueVoidPtr.h
```

```
class UniqueVoidPtr {
private:
    // Lifetime tied to ctx_
    void* data_;
    std::unique_ptr<void, DeleterFnPtr> ctx_;

    // ...
}
```

现在我们知道，在C++的层面，张量被Tensor类型所表示，但是我们平时是使用Python语言来训练推理模型的，使用详细的过程我们留到后面的章节解释，不过机制并不复杂，PyTorch使用了THPVariable这个类型作为过渡，Python中在前面初始化_C模块的时候，调用了THPVariable_initModule()这个函数，将Python中_TensorBase这个类型映射到

```
// torch/csrc/autograd/python_variable.cpp
```

```
bool THPVariable_initModule(PyObject *module)
{
    // ...
    PyModule_AddObject(module, "_TensorBase", (PyObject *)&THPVariableType);
    torch::autograd::initTorchFunctions(module);
    // ...
    return true;
}
```

```
PyTypeObject THPVariableType = {
    PyVarObject_HEAD_INIT(
        &THPVariableMetaType,
        0) "torch._C._TensorBase", /* tp_name */
    // ...
    THPVariable_pynew, /* tp_new */
};
```

```
PyObject *THPVariable_pynew(PyTypeObject *type, PyObject *args, PyObject *kwargs)
{
    HANDLE_TH_ERRORS
    TORCH_CHECK(type != &THPVariableType, "Cannot directly construct _TensorBase; subclass it");
    jit::tracer::warn("torch.Tensor", jit::tracer::WARN_CONSTRUCTOR);
    auto tensor = torch::utils::base_tensor_ctor(args, kwargs);
    // WARNING: tensor is NOT guaranteed to be a fresh tensor; e.g., if it was
    // given a raw pointer that will refcount bump
    return THPVariable_NewWithVar(
        type,
        std::move(tensor),
```

```

        c10::impl::PyInterpreterStatus::MAYBE_UNINITIALIZED);
    END_HANDLE_TH_ERRORS
}

static PyObject* THPVariable_NewWithVar(
    PyTypeObject* type,
    Variable _var,
    c10::impl::PyInterpreterStatus status) {

    PyObject* obj = type->tp_alloc(type, 0);
    if (obj) {
        auto v = (THPVariable*) obj;
        // TODO: named constructor to avoid default initialization
        new (&v->cdata) MaybeOwned<Variable>();
        v->cdata = MaybeOwned<Variable>::owned(std::move(_var));
        const auto& var = THPVariable_Unpack(v);
        var.unsafeGetTensorImpl()->init_pyobj(self_interpreter.get(), obj, status);
        if (check_has_torch_dispatch(obj)) {
            var.unsafeGetTensorImpl()->set_python_dispatch(true);
        }
    }
    return obj;
}

// torch/csrc/autograd/python_variable.h
struct THPVariable {
    PyObject_HEAD;
    c10::MaybeOwned<at::Tensor> cdata;
    PyObject* backward_hooks = nullptr;
};

```

TensorOption

Note: 参考注释吧

TensorOption是设计用来构造Tensor的工具。

在C++中没有python中的keyword参数机制，比如这段代码：

```
torch.zeros(2, 3, dtype=torch.int32)
```

在keyword参数机制下，参数的顺序和定义的可能不一样。因此在C++中实现这些函数时，将TensorOptions作为最后

实际使用时，at::zeros()系列函数隐式的使用TensorOptions。 TensorOptions
s可以看作是一个字典。

Node

Node的定义在torch/csrc/autograd/function.h中。

从名称上不难看出，Node代表计算图中的节点。计算图除了节点之外，还会有边，也就是Edge。

Tensor中方法grad_fn()返回的就是一个Node

Edge

Node的定义在torch/csrc/autograd/edge.h中。

VariableHooks

获取Tensor的grad_fn()时，使用VariableHooks这个类来返回的，而且逻辑很复杂，还没看懂

<https://blog.csdn.net/u012436149/article/details/69230136>

这里要注意的是，hook 只能注册到 Module 上，即，仅仅是简单的 op 包装的 Module，而不是我们继承 Module时写的那个类，我们继承 Module写的类叫做 Container。每次调用forward()计算输出的时候，这个hook就会被调用。它应该拥有以下签名：

可以看到，当我们执行model(x)的时候，底层干了以下几件事：

```
forward
    forward_hook    forward    hook    hook
```

register_backward_hook

在module上注册一个backward hook。此方法目前只能用在Module上，不能用在Container上，当Module的forward函数每次计算module的inputs的梯度的时候，这个hook会被调用。hook应该拥有下面的signature。

hook(module, grad_input, grad_output) -> Tensor or None

如果module有多个输入输出的话，那么grad_input grad_output将会是个tuple。

hook不应该修改它的arguments，但是它可以选择性的返回关于输入的梯度，这个返回的梯度在后续的计算中会替代这个函数返回一个句柄(handle)。它有一个方法 handle.remove()，可以用这个方法将hook从module移除。

从上边描述来看，backward hook似乎可以帮助我们处理一下计算完的梯度。看下面nn.Module中register_backward_hook

Backward函数注册流程

```
initialize_autogenerated_functionsEverything();
addClass<AddBackward0>(AddBackward0Class,"AddBackward0", AddBackward0_properties);
    _initFunctionPyTypeObject();
```

```
registerCppFunction();
    cpp_function_types[idx] = type
```

参考

- <https://blog.csdn.net/Xixo0628/article/details/112603174>
- <https://blog.csdn.net/Xixo0628/article/details/112603174>
- <https://pytorch.org/blog/a-tour-of-pytorch-internals-1/#the-thptensor-type>
- PyTorch源码浅析(1): THTensor <https://blog.csdn.net/Xixo0628/article/details/112603174>
- PyTorch源码浅析(1): THTensor <https://www.52coding.com.cn/2019/05/05/PyTorch1/>

PyTorch的算子体系

主要内容

- PyTorch中算子的实现方式
- 源代码的组织
- 运行代码分析
- 自定义算子的实现
- torch模块中的函数
- Tensor算子
- torch.nn的算子
- 算子的注册过程
- 算子的调用过程

一个简单的例子

我们先从一个简单的例子出发，看看PyTorch中Python和C++是怎样一起工作的。

```
import torch
```

```
x = torch.ones(2, 2, requires_grad=True)
y = x + 2
```

在_C模块初始化的时候，THPVariable这个类型绑定了相应的方法，可以在执行加法操作的时候，调用的是THPVariable

```
PyMethodDef variable_methods[] = {
    // These magic methods are all implemented on python object to wrap NotImplementedError
    {"__add__", castPyCFunctionWithKeywords(TypeError_to_NotImplemented_<THPVariable_add>), METH_VARARGS},
    {"__radd__", castPyCFunctionWithKeywords(TypeError_to_NotImplemented_<THPVariable_add>), METH_VARARGS},
    {"__iadd__", castPyCFunctionWithKeywords(TypeError_to_NotImplemented_<THPVariable_add>), METH_VARARGS},
}
```

```
...
}
```

THPVariable_add() 方法的具体实现代码是生成的，因此我们在原始的模板文件中可以找到使用这个函数，真正的实现

```
// torch/csrc/autograd/generated/python_variable_methods.cpp [generated file]
```

```
static PyObject * THPVariable_add(PyObject* self_, PyObject* args, PyObject* kwargs)
{
    HANDLE_TH_ERRORS
    const Tensor& self = THPVariable_Unpack(self_);
    static PythonArgParser parser({
        "add(Scalar alpha, Tensor other)|deprecated",
        "add(Tensor other, *, Scalar alpha=1)",
    }, /*traceable=*/true);

    ParsedArgs<2> parsed_args;
    auto _r = parser.parse(self_, args, kwargs, parsed_args);
    if(_r.has_torch_function()) {
        return handle_torch_function(_r, self_, args, kwargs, THPVariableClass, "torch.Tensor");
    }
    switch (_r.idx) {
    case 0: {
        // [deprecated] aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor

        auto dispatch_add = [](const at::Tensor & self, const at::Scalar & alpha, const at::Tensor & other) {
            pybind11::gil_scoped_release no_gil;
            return self.add(other, alpha);
        };
        return wrap(dispatch_add(self, _r.scalar(0), _r.tensor(1)));
    }
    case 1: {
        // aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor

        auto dispatch_add = [](const at::Tensor & self, const at::Tensor & other, const at::Scalar & alpha) {
            pybind11::gil_scoped_release no_gil;
            return self.add(other, alpha);
        };
        return wrap(dispatch_add(self, _r.tensor(0), _r.scalar(1)));
    }
    }
    Py_RETURN_NONE;
    END_HANDLE_TH_ERRORS
}
```

其中 PythonArgParser 定义了这个函数的几类参数，并将Python调用的参数转换成对应的C++类型，在这个例子里，

```
// aten/src/ATen/core/TensorBody.h
```

```

// aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor
inline at::Tensor Tensor::add(const at::Tensor & other, const at::Scalar & alpha) const {
    return at::_ops::add_Tensor::call(const_cast<Tensor&>(*this), other, alpha);
}

// ./build/aten/src/ATen/Operators_2.cpp [generated file]

STATIC_CONST_STR_OUT_OF_LINE_FOR_WIN_CUDA(add_Tensor, name, "aten::add")
STATIC_CONST_STR_OUT_OF_LINE_FOR_WIN_CUDA(add_Tensor, overload_name, "Tensor")
STATIC_CONST_STR_OUT_OF_LINE_FOR_WIN_CUDA(add_Tensor, schema_str, "add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor")

// aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor
static C10_NOINLINE c10::TypedOperatorHandle<add_Tensor::schema> create_add_Tensor_typed_handle() {
    return c10::Dispatcher::singleton()
        .findSchemaOrThrow(add_Tensor::name, add_Tensor::overload_name)
        .typed<add_Tensor::schema>();
}

// aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor
at::Tensor add_Tensor::call(const at::Tensor & self, const at::Tensor & other, const at::Scalar & alpha) {
    static auto op = create_add_Tensor_typed_handle();
    return op.call(self, other, alpha);
}

```

这里创建的op的类型是c10::OperatorHandle

算子分发的基本概念

增加新的算子时，需要先使用TORCH_LIBRARY定义算子的schema，然后使用宏TORCH_LIBRARY_IMPL来注册该算子在cpu、cuda、XLA等上的实现。注册的时候，需要指定namespace及该namespace下的实现。

参考官方文档 <https://pytorch.org/tutorials/advanced/dispatcher.html>

在了解Dispatch的机制之前，我们先了解一下算子的类型。

BackendComponent

每一种“backend”可以看做是一种设备。

```
// c10/core/DispatchKey.h
```

```
enum class BackendComponent : uint8_t {
    InvalidBit = 0,
    CPUBit,
    ...
}
```

```

    CUDABit,
    HIPBit,
    XLABit,
    MPSBit,
    IPUBit,
    XPUBit,
    HPUBit,
    VEBit,
    LazyBit,
    PrivateUse1Bit,
    PrivateUse2Bit,
    PrivateUse3Bit,
    // Define an alias to represent end of backend dispatch keys.
    // If you add new backend keys after PrivateUse3, please also update it here.
    // (But you shouldn't: private use keys should have higher precedence than
    // all built-in keys)
    EndOfBackendKeys = PrivateUse3Bit,
};

DispatchKey

// c10/core/DispatchKey.h

enum class DispatchKey : uint16_t {

    Undefined = 0,

    CatchAll = Undefined,

    // ~~~~~~ Functionality Keys ~~~~~~ //
    Dense,

    // Below are non-extensible backends.
    // These are backends that currently don't have their own overrides for
    // Autograd/Sparse/Quantized kernels,
    // and we therefore don't waste space in the runtime operator table allocating
    // space for them.
    // If any of these backends ever need to customize, e.g., Autograd, then we'll
    // need to add a DispatchKey::*Bit for them.

    FPGA, // Xilinx support lives out of tree at
    // https://gitlab.com/pytorch-complex/vitis\_kernels

    // ONNX Runtime, lives out of tree at https://github.com/pytorch/ort and
    // https://github.com/microsoft/onnxruntime, and is also used to test general

```



```

// backend/extension machinery in the core. cf:
// - test/cpp_extensions/ort_extension.cpp
// - test/test_torch.py
// - aten/src/ATen/test/extension_backend_test.cpp
ORT,

Vulkan,
Metal,

// A meta tensor is a tensor without any data associated with it. (They
// have also colloquially been referred to as tensors on the "null" device).
// A meta tensor can be used to dry run operators without actually doing any
// computation, e.g., add on two meta tensors would give you another meta
// tensor with the output shape and dtype, but wouldn't actually add anything.
Meta,

// See [Note: Per-Backend Functionality Dispatch Keys]
Quantized,

// This backend is to support custom RNGs; it lets you go
// to a different kernel if you pass in a generator that is not a
// traditional CPUGeneratorImpl/CUDAGeneratorImpl. To make use of this
// key:
// 1) set it as a second parameter of at::Generator constructor call in
// the user-defined PRNG class.
// 2) use it as a dispatch key while registering custom kernels
// (templated kernels specialized for user-defined PRNG class)
// intended for out of tree use; tested by aten/src/ATen/test/rng_test.cpp
CustomRNGKeyId,

// Here are backends which specify more specialized operators
// based on the layout of the tensor. Note that the sparse backends
// are one case where ordering matters: sparse multi-dispatches with
// the corresponding dense tensors, and must be handled before them.
MkldnnCPU, // registered at build/aten/src/ATen/RegisterMkldnnCPU.cpp
// NB: not to be confused with MKLDNN, which is Caffe2 only

// See [Note: Per-Backend Functionality Dispatch Keys]
Sparse,

SparseCsrCPU,
SparseCsrCUDA,

// Note [Non-Customizable Backend Keys]
// Every key above here is considered a "non-customizable backend".
// These are backends that will work correctly with autograd, but

```

```

// but currently don't require separate implementations
// for autograd sparse or quantized kernels.
// Any new backends that don't need to be customized should go above here.
// If an existing backend needs to e.g. override autograd, then we can
// consider promoting it into the "BackendComponent" enum
//
// For all intents and purposes from the perspective of DispatchKeySet,
// "non-customizable backend" keys are treated the same way
// as other functionality keys
EndOfNonCustomizableBackends = SparseCsrCUDA,

NestedTensor,

// In some situations, it is not immediately obvious what the correct
// backend for function is, because the function in question doesn't
// have any "tensor" arguments. In this case, a BackendSelect function
// can be registered to implement the custom determination of the
// correct backend.
BackendSelect,

Python,

// Out-of-core key for Fake Tensor in torchdistx.
// See https://pytorch.org/torchdistx/latest/fake\_tensor.html
Fake,

// The named dispatch key is set for any tensors with named dimensions.
// Although we have a dispatch key for named tensors, for historical reasons,
// this dispatch key doesn't do any of the substantive functionality for named
// tensor (though, hypothetically, it could!) At the moment, it's just
// responsible for letting us give good error messages when operations
// don't support named tensors.
//
// NB: If you ever consider moving named tensor functionality into
// this dispatch key, note that it might be necessary add another dispatch
// key that triggers before composite operators, in case a composite operator
// has named dimension propagation that doesn't match that of its
// constituent parts.
Named,

// The Conjugate dispatch key is set for any tensors that need to perform
// conjugation
// This is implemented at a dispatch level right before any backends run
Conjugate,

// The Negative dispatch key is set for any tensors that need to perform

```

```

// negation
// This is implemented at a dispatch level right before any backends run
Negative,

ZeroTensor, // registered at build/aten/src/ATen/RegisterZeroTensor.cpp

// See Note [Out-of-tree vmap+grad prototype]. The purpose of this key
// is to insert code after the "autograd subsystem" runs, so this key should
// be directly after ADInplaceOrView and all of the autograd keys.
FuncTorchDynamicLayerBackMode,

// Note [ADInplaceOrView key]
// ADInplaceOrView key is used by inplace or view ops to register a kernel
// that does additional setup for future autograd computation.
//
// 1. For inplace ops this kernel does version bump
// 2. For view ops this kernel does `as_view` setup where we properly setup
//    DifferentiableViewMeta on the view tensors.
//
// For other ops it's fallthrough kernel since there's no extra
// work to do.
//
// Note [Dream: skip VariableType kernel when requires_grad=false]
//
// In an ideal world where we can skip VariableType kernel for inputs
// with requires_grad=false, instead of a fallthrough kernel, we'll
// register a kernel shown below to all functional ops as well:
// torch::Tensor my_functional_op(...) {
//     {
//         // Note for every op in VariableType, you need to go through
//         // `AutoDispatchBelowADInplaceOrView` guard exactly once to add the
//         // key to TLS excluded set. If you don't go through it at all,
//         // inplace/view ops called through `at::` inside your backend
//         // kernel will dispatch to ADInplaceOrView kernels and do a lot
//         // of extra work.
//         at::AutoDispatchBelowADInplaceOrView guard;
//         at::redispatch::my_functional_op(...);
//     }
// }
// But this work is currently blocked since it adds an extra dispatch
// for all ops and it's non-trivial overhead at model level(a few percents).
// Thus our current approach takes advantage of the fact every kernel go
// through VariableType kernel first and pulls the
// `at::AutoDispatchBelowADInplaceOrView` guard of functional ops
// up to the `VariableType` kernel. Thus we only add the extra dispatch
// to view/inplace ops to minimize its perf impact to real models.

```

```

ADInplaceOrView,
// Note [Alias Dispatch Key : Autograd]
// All backends are oblivious to autograd; autograd is handled as a
// layer which happens on top of all backends. It inspects the autograd
// metadata of all inputs, determines what autograd metadata should be
// constructed by the output, and otherwise defers to the backend to
// actually do the numeric computation. Autograd contains
// the bulk of this logic.

// Autograd is now an alias dispatch key which by default maps to all
// backend-specific autograd keys.
// Backend-specific allow backends to override the default kernel registered
// to Autograd key as needed.
// For example, XLA wants to define autograd for einsum directly.
// Registering a custom autograd implementation at the XLA key won't work
// because we process Autograd before XLA. This key has higher priority and
// gets processed first. You generally should NOT redispach after handling
// autograd here (since that would result in execution of the Autograd
// operator, which you're trying to skip). In AutogradXLA implementations,
// you are responsible for handling autograd yourself, or deferring to other
// operators which support autograd.

// Currently we only have backend-specific autograd keys for CPU/CUDA/XLA and
// reserved user-defined backends. All other in-tree backends share the
// AutogradOther key. We can add specific autograd key for those backends
// upon request.
AutogradOther,

// See [Note: Per-Backend Functionality Dispatch Keys]
AutogradFunctionality,

// NestedTensor is an example of something that isn't a "real backend"
// (because it mostly consists of redispaching kernels)
// but it would like to override autograd functionality in C++.
// We can handle cases like this by adding an extra functionality key
// exclusively for handling autograd for NestedTensor.
// lives out of tree at
// https://github.com/pytorch/nestedtensor
AutogradNestedTensor,

Tracer,

// Autocasting precedes VariableTypeId, to ensure casts are autograd-exposed
// and inputs are saved for backward in the post-autocast type.
AutocastCPU,
AutocastXPU,

```

```

// Naughtily, AutocastCUDA is also being used for XLA. In the terminal state,
// it probably should get its own Autocast key
AutocastCUDA,

// ~~~~~ WRAPPERS ~~~~~ //
// There are a number of alternative modes which may want to handle before
// autograd; for example, error checking, tracing, profiling or vmap. They
// go here.

FuncTorchBatched, // See Note [Out-of-tree vmap+grad prototype]
FuncTorchVmapMode, // See Note [Out-of-tree vmap+grad prototype]

// This is the dispatch key for BatchedTensorImpl, which is used to implement
// batching rules for vmap.
Batched,

// When we are inside a vmap, all tensors dispatch on this key.
// See Note: [DispatchKey::VmapMode usage] for more details.
VmapMode,

FuncTorchGradWrapper, // See Note [Out-of-tree vmap+grad prototype]

// Alias and mutation removal.
// If some backends want to opt into only alias removal or only mutation
// removal,
// we can consider adding separate keys dedicated to those individual passes.
// See Note [Functionalization Pass In Core] for details.
Functionalize,

// Out-of-core key for Deferred Module Initialization in torchdistx.
// See https://pytorch.org/torchdistx/latest/deferred\_init.html
DeferredInit,

// Used by Python key logic to know the set of tls on entry to the dispatcher
// This kernel assumes it is the top-most non-functorch-related DispatchKey.
// If you add a key above, make sure to update the fallback implementation for
// this.
PythonTLSSnapshot,

// This key should be at the very top of the dispatcher
FuncTorchDynamicLayerFrontMode, // See Note [Out-of-tree vmap+grad prototype]

// TESTING: This is intended to be a generic testing tensor type id.
// Don't use it for anything real; its only acceptable use is within a single
// process test. Use it by creating a TensorImpl with this DispatchKey, and
// then registering operators to operate on this type id. See

```

```

// aten/src/ATen/core/dispatch/backend_fallback_test.cpp for a usage example.
TESTING_ONLY_GenericWrapper,

// TESTING: This is intended to be a generic testing tensor type id.
// Don't use it for anything real; its only acceptable use is within a single
// process test. Use it by toggling the mode on and off via
// TESTING_ONLY_tls_generic_mode_set_enabled and then registering operators
// to operate on this type id. See
// aten/src/ATen/core/dispatch/backend_fallback_test.cpp
// for a usage example
TESTING_ONLY_GenericMode,

// ~~~~~ FIN ~~~~~ //
EndOfFunctionalityKeys, // End of functionality keys.

// ~~~~~ "Dense" Per-Backend Dispatch keys ~~~~~ //
// Here are backends which you think of as traditionally specifying
// how to implement operations on some device.

// See Note [The Ordering of Per-Backend Dispatch Keys Matters!]
StartOfDenseBackends,
CPU, // registered at build/aten/src/ATen/RegisterCPU.cpp
CUDA, // registered at build/aten/src/ATen/RegisterCUDA.cpp
HIP, // NB: I think this is not actually used, due to Note [Masquerading as
// CUDA]
XLA, // lives out of tree at https://github.com/pytorch/xla
MPS, // registered at build/aten/src/ATen/RegisterMPS.cpp
IPU, // lives out of tree at https://github.com/graphcore/poptorch
XPU, // For out of tree Intel's heterogeneous computing plug-in
HPU, // For out of tree & closed source integration of HPU / Habana
VE, // For out of tree & closed source integration of SX-Aurora / NEC
Lazy, // For lazy tensor backends
// Here are reserved backends for user-defined backends, see Note [Private use
// DispatchKey]
// To see some example about how to use this, check out ORT
PrivateUse1,
PrivateUse2,
PrivateUse3,
EndOfDenseBackends = PrivateUse3,

// ~~~~~ "Quantized" Per-Backend Dispatch keys ~~~~~ //
// keys starting with an _ are not currently used,
// but are needed to ensure that every backend is indexed correctly.

// See Note [The Ordering of Per-Backend Dispatch Keys Matters!]
StartOfQuantizedBackends,

```

```

QuantizedCPU, // registered at build/aten/src/ATen/RegisterQuantizedCPU.cpp
QuantizedCUDA, // registered at build/aten/src/ATen/RegisterQuantizedCUDA.cpp
_QuantizedHIP,
_QuantizedXLA,
_QuantizedMPS,
_QuantizedIPU,
QuantizedXPU, // For out of tree Intel's heterogeneous computing plug-in
_QuantizedHPU,
_QuantizedVE,
_QuantizedLazy,
_QuantizedPrivateUse1,
_QuantizedPrivateUse2,
_QuantizedPrivateUse3,
EndOfQuantizedBackends = _QuantizedPrivateUse3,

// ~~~~~ "Sparse" Per-Backend Dispatch keys ~~~~~ //
// keys starting with an _ are not currently used,
// but are needed to ensure that every backend is indexed correctly.

// See Note [The Ordering of Per-Backend Dispatch Keys Matters!]
StartOfSparseBackends,
SparseCPU, // registered at build/aten/src/ATen/RegisterSparseCPU.cpp
SparseCUDA, // registered at build/aten/src/ATen/RegisterSparseCUDA.cpp
SparseHIP, // TODO: I think this is not actually used, due to Note
// [Masquerading as CUDA]
_SparseXLA,
_SparseMPS,
_SparseIPU,
SparseXPU, // For out of tree Intel's heterogeneous computing plug-in
_SparseHPU,
SparseVE, // For out of tree & closed source integration of SX-Aurora / NEC
_SparseLazy,
_SparsePrivateUse1,
_SparsePrivateUse2,
_SparsePrivateUse3,
EndOfSparseBackends = _SparsePrivateUse3,

// ~~~~~ "NestedTensor" Per-Backend Dispatch keys ~~~~~
// //
// keys starting with an _ are not currently used,
// but are needed to ensure that every backend is indexed correctly.

// See Note [The Ordering of Per-Backend Dispatch Keys Matters!]
StartOfNestedTensorBackends,
// registered at build/aten/src/ATen/RegisterNestedTensorCPU.cpp
NestedTensorCPU,

```

```

// registered at build/aten/src/ATen/RegisterNestedTensorCUDA.cpp
NestedTensorCUDA,
_NestedTensorHIP,
_NestedTensorXLA,
_NestedTensorMPS,
_NestedTensorIPU,
_NestedTensorXPU,
_NestedTensorHPU,
_NestedTensorVE,
_NestedTensorLazy,
_NestedTensorPrivateUse1,
_NestedTensorPrivateUse2,
_NestedTensorPrivateUse3,
EndOfNestedTensorBackends = _NestedTensorPrivateUse3,

// ~~~~~ "Autograd" Per-Backend Dispatch keys ~~~~~ //
// keys starting with an _ are not currently used,
// but are needed to ensure that every backend is indexed correctly.

// See Note [The Ordering of Per-Backend Dispatch Keys Matters!]
StartOfAutogradBackends,
AutogradCPU,
AutogradCUDA,
_AutogradHIP,
AutogradXLA,
AutogradMPS,
AutogradIPU,
AutogradXPU,
AutogradHPU,
_AutogradVE,
AutogradLazy,
// Here are some reserved pre-autograd keys for user-defined backends, see
// Note [Private use DispatchKey]
AutogradPrivateUse1,
AutogradPrivateUse2,
AutogradPrivateUse3,
EndOfAutogradBackends = AutogradPrivateUse3,
// If we add a new per-backend functionality key that has higher priority
// than Autograd, then this key should be updated.
EndOfRuntimeBackendKeys = EndOfAutogradBackends,

// ~~~~~ Alias Dispatch Keys ~~~~~ //
// Note [Alias Dispatch Keys]
// Alias dispatch keys are synthetic dispatch keys which map to multiple
// runtime dispatch keys. Alias keys have precedence, but they are always
// lower precedence than runtime keys. You can register a kernel to an

```



```

// alias key, the kernel might be populated to the mapped runtime keys
// during dispatch table computation.
// If a runtime dispatch key has multiple kernels from alias keys, which
// kernel wins is done based on the precedence of alias keys (but runtime
// keys always have precedence over alias keys).
// Alias keys won't be directly called during runtime.

// See Note [Alias Dispatch Key : Autograd]
Autograd,
CompositeImplicitAutograd, // registered at
// build/aten/src/ATen/RegisterCompositeImplicitAutograd.cpp
CompositeExplicitAutograd, // registered at
// build/aten/src/ATen/RegisterCompositeExplicitAutograd.cpp

// Define an alias key to represent end of alias dispatch keys.
// If you add new alias keys after Autograd, please also update it here.
StartOfAliasKeys = Autograd,
EndOfAliasKeys = CompositeExplicitAutograd, //

// ~~~~~ BC ALIASES ~~~~~ //
// The aliases exist for backwards compatibility reasons, they shouldn't
// be used
CPUTensorId = CPU,
CUDATensorId = CUDA,
DefaultBackend = CompositeExplicitAutograd,
PrivateUse1_PreAutograd = AutogradPrivateUse1,
PrivateUse2_PreAutograd = AutogradPrivateUse2,
PrivateUse3_PreAutograd = AutogradPrivateUse3,
Autocast = AutocastCUDA,
};

```

DispatchKeySet

所有的算子都是注册在Dispatcher里的，在调用的时候，根据函数名词和传递的参数类型，dispatcher会寻找相应的

下面内容来自PyTorch源码中对DispatchKeySet的注释（翻译不准确的请指正）：>

DispatchKeySet就是一组DispatchKey，包括了“functionality”和“backend”两种比特位，每个tensor都有自己

> Dispatcher根据tensor的keyset或者多个tensor的keyset组合，实现了不同的dispatch，并分发到不同的实现（f

> 在内部实现上，Dispatch key 被打包成64位的DispatchKeySet对象。>

总的key的数量是[backends] * [functionalities]，因此直接把每个key与每个bit关联是不太合适的，key太多了，

> 两个枚举值（BackendComponent和DispatchKey）可以被分为5个类别：> (1)

“Building block” keys > (a) backends: BackendComponent枚举，比如CPUBit,

CUDABit > (b) functionalities (per-backend) 功能相关的dispatch key, 比如

AutogradFunctionality, Sparse, Dense > (2) “Runtime” keys > (a)

“non-customizable backends”，比如FPGA > (b) “non-customizable function-

alites”，比如Functionalize > (c) “per-backend instances of customizable functionalities”，比如CPU, SparseCPU, AutogradCPU > (3) “Alias” DispatchKeys > > (1) Building block的key可以组合成一个运行时使用的DispatchKeySet，例如：

```
> auto dense_cpu_ks = DispatchKeySet({DispatchKey::CPUBit, > DispatchKey::Dense}); > // The keyset has the runtime dense-cpu key.
> dense_cpu_ks.has(DispatchKey::CPU); > // And it contains the building block keys too. > dense_cpu_ks.has(DispatchKey::CPUBit); > dense_cpu_ks.has(DispatchKey::Dense); > 但不是所有的backend或者functionality都可以作为building block，这样就允许了更灵活的设计 > #### Dispatcher
```

Dispatcher的作用是根据实际的上下文选择不同的operator实现，

```
class TORCH_API Dispatcher final {
private:

    struct OperatorDef final { ... };

public:
    static Dispatcher& realSingleton();

    C10_ALWAYS_INLINE static Dispatcher& singleton() { ... }

    c10::optional<OperatorHandle> findSchema(const OperatorName& operator_name);

    OperatorHandle findSchemaOrThrow(const char* name, const char* overload_name);

    c10::optional<OperatorHandle> findOp(const OperatorName& operator_name);

    const std::vector<OperatorName> getAllOpNames();

    template<class Return, class... Args>
    Return call(const TypedOperatorHandle<Return (Args...)>& op, Args... args) const;

    template<class Return, class... Args>
    Return redispatch(const TypedOperatorHandle<Return (Args...)>& op, DispatchKeySet currentKeys) const;

    // Invoke an operator via the boxed calling convention using an IValue stack
    void callBoxed(const OperatorHandle& op, Stack* stack) const;

    // TODO: This will only be useful if we write a backend fallback that plumbs dispatch keys
    // See Note [Plumbing Keys Through The Dispatcher]
    void redispatchBoxed(const OperatorHandle& op, DispatchKeySet dispatchKeySet, Stack* stack) const;

    RegistrationHandleRAII registerDef(FunctionSchema schema, std::string debug);
    RegistrationHandleRAII registerImpl(OperatorName op_name, c10::optional<DispatchKey> dispatchKey);
};
```

```

RegistrationHandlerRAII registerName(OperatorName op_name);

RegistrationHandlerRAII registerFallback(DispatchKey dispatch_key, KernelFunction kernel, s

RegistrationHandlerRAII registerLibrary(std::string ns, std::string debug);

std::vector<OperatorName> getRegistrationsForDispatchKey(c10::optional<DispatchKey> k) con

private:
    // ...

    std::list<OperatorDef> operators_;
    LeftRight<ska::flat_hash_map<OperatorName, OperatorHandle>> operatorLookupTable_;
    ska::flat_hash_map<std::string, std::string> libraries_;

    std::array<impl::AnnotatedKernel, num_runtime_entries> backendFallbackKernels_;

    // ...
};

```

算子注册过程

在PyTorch中，全局只有一个唯一的Dispatcher，所有的算子都注册到这个Dispatcher上，因为算子很多，为了方便

TORCH_LIBRARY及Schema说明

TORCH_LIBRARY可以用来注册Schema，在aten这个namespace下，就注册了超过2500个schema。

```
// build/aten/src/ATen/RegisterSchema.cpp
```

```

TORCH_LIBRARY(aten, m) {
    // ...
    m.def("cudnn_batch_norm(Tensor input, Tensor weight, Tensor? bias, Tensor? running_mean, T
    m.def("cudnn_batch_norm_backward(Tensor input, Tensor grad_output, Tensor weight, Tensor?
    m.def("cudnn_convolution(Tensor self, Tensor weight, int[] padding, int[] stride, int[] di
    m.def("cudnn_convolution_transpose(Tensor self, Tensor weight, int[] padding, int[] output
        // ...
}

```

我们看一下 TORCH_LIBRARY被定义在torch/library.h中，从这个文件的位置也可以看出其重要性。这个宏有两个参Library.

```
// torch/library.h
```

```
class TorchLibraryInit final {
```

```

private:
    using InitFn = void(Library&);
    Library lib_;

public:
    TorchLibraryInit(
        Library::Kind kind,
        InitFn* fn,
        const char* ns,
        c10::optional<c10::DispatchKey> k,
        const char* file,
        uint32_t line)
        : lib_(kind, ns, k, file, line) {
        fn(lib_);
    }
};

```

```

#define TORCH_LIBRARY(ns, m) \
    static void TORCH_LIBRARY_init_##ns(torch::Library&); \
    static const torch::detail::TorchLibraryInit TORCH_LIBRARY_static_init_##ns( \
        torch::Library::DEF, \
        &TORCH_LIBRARY_init_##ns, \
        #ns, \
        c10::nullopt, \
        __FILE__, \
        __LINE__); \
    void TORCH_LIBRARY_init_##ns(torch::Library& m)

```

在这个宏里，首先声明一个算子库的初始化函数，然后创建了一个TorchLibraryInit的实例，这个实例会初始化Lib。在Library的实例化过程中，该Library也会被注册到全局的Dispatcher里，如下面的实现所示，注册的时候以names

// aten/src/ATen/core/library.cpp

```

Library::Library(Kind kind, std::string ns, c10::optional<c10::DispatchKey> k, const char* file, uint32_t line)
    : kind_(kind)
    , ns_(ns == "_" ? c10::nullopt : c10::make_optional(std::move(ns)))
    , dispatch_key_((!k.has_value() || *k == c10::DispatchKey::CatchAll) ? c10::nullopt : k)
    , file_(file)
    , line_(line)
    {
    switch (kind_) {
    case DEF:
        registrars_.emplace_back(
            c10::Dispatcher::singleton().registerLibrary(
                *ns_, debugString(file_, line_)
            )
        );
    };
}

```

```

    case FRAGMENT:
        //...
        break;
    case IMPL:
        // Nothing to do, everything is OK
        break;
}
}

```

TODO: add schema specification

TORCH_LIBRARY_IMPL

每个算子有唯一的schema，但是可能有很多的实现，在实际运行中，PyTorch会通过Dispatcher查找合适的实现并执行。算子实现的注册方式是通过TORCH_LIBRARY_IMPL，例如，在下面的代码中，注册了多个Autograd算子和CUDA。

// torch/csrc/autograd/generated/VariableTypeEverything.cpp

```

TORCH_LIBRARY_IMPL(aten, Autograd, m) {
    // ...
    m.impl("add.Tensor",
           TORCH_FN(VariableType::add_Tensor)
    );
    m.impl("add.Scalar",
           TORCH_FN(VariableType::add_Scalar)
    );
    // ...
}

```

// build/aten/src/ATen/RegisterCPU.cpp

```

TORCH_LIBRARY_IMPL(aten, CPU, m) {
    // ...
    m.impl("add.Tensor", TORCH_FN(wrapper_add_Tensor));
    m.impl("add.out", TORCH_FN(wrapper_add_out_out));
    // ...
}

```

// build/aten/src/ATen/RegisterCUDA.cpp

```

TORCH_LIBRARY_IMPL(aten, CUDA, m) {
    //...
    m.impl("cudnn_batch_norm",
           TORCH_FN(wrapper__cudnn_batch_norm));

    m.impl("cudnn_batch_norm_backward",
           TORCH_FN(wrapper__cudnn_batch_norm_backward));
}

```

```

    m.impl("cudnn_convolution",
    TORCH_FN(wrapper__cudnn_convolution));

    m.impl("cudnn_convolution_transpose",
    TORCH_FN(wrapper__cudnn_convolution_transpose));
    //...
}

```

容易看出，TORCH_LIBRARY_IMPL定义了命名空间ns下，DispatchKeySet为CUDA的一组算子实现，开发者可以通过m.i
下面我们看一下这个宏的实现：

```

// torch/library.h

#define TORCH_LIBRARY_IMPL(ns, k, m) _TORCH_LIBRARY_IMPL(ns, k, m, C10_UID)

#define _TORCH_LIBRARY_IMPL(ns, k, m, uid) \
    static void C10_CONCATENATE( \
        TORCH_LIBRARY_IMPL_init_##ns##_##k##_ , uid)(torch::Library&); \
    static const torch::detail::TorchLibraryInit C10_CONCATENATE( \
        TORCH_LIBRARY_IMPL_static_init_##ns##_##k##_ , uid)( \
        torch::Library::IMPL, \
        c10::guts::if_constexpr<c10::impl::dispatch_key_allowlist_check( \
            c10::DispatchKey::k)>( \
            []() { \
                return &C10_CONCATENATE( \
                    TORCH_LIBRARY_IMPL_init_##ns##_##k##_ , uid); \
            }, \
            []() { return [] (torch::Library&) -> void {}; } ), \
        #ns, \
        c10::make_optional(c10::DispatchKey::k), \
        __FILE__, \
        __LINE__); \
    void C10_CONCATENATE( \
        TORCH_LIBRARY_IMPL_init_##ns##_##k##_ , uid)(torch::Library & m)

```

和宏TORCH_LIBRARY类似，TORCH_LIBRARY_IMPL首先声明一个算子库的初始化函数，然后创建了一个TorchLibraryIn
在Library的实例化过程中，该Library也会被注册到全局的Dispatcher里，如下面的实现所示，注册的时候以names

接下来我们看一下注册方法实现的细节，因为算子对应的实现，也就是kernel
function，是通过m.impl()来注册的，我们看一下该方法的实现：

```

// aten/src/ATen/core/library.cpp

Library& Library::_impl(const char* name_str, CppFunction&& f) & {
    auto name = torch::jit::parseName(name_str);
    auto ns_opt = name.getNamespace();

```

```

//...

auto dispatch_key = f.dispatch_key_.has_value() ? f.dispatch_key_ : dispatch_key_;
registrars_.emplace_back(
    c10::Dispatcher::singleton().registerImpl(
        std::move(name),
        dispatch_key,
        std::move(f.func_),
        // NOLINTNEXTLINE(performance-move-const-arg)
        std::move(f.cpp_signature_),
        std::move(f.schema_),
        debugString(std::move(f.debug_), file_, line_)
    )
);
return *this;
}

// aten/src/ATen/core/dispatch/Dispatcher.cpp
RegistrationHandlerRAII Dispatcher::registerImpl(
    OperatorName op_name,
    c10::optional<DispatchKey> dispatch_key,
    KernelFunction kernel,
    c10::optional<impl::CppSignature> cpp_signature,
    std::unique_ptr<FunctionSchema> inferred_function_schema,
    std::string debug
) {
    std::lock_guard<std::mutex> lock(mutex_);

    auto op = findOrRegisterName_(op_name);

    auto handle = op.operatorDef_->op.registerKernel(
        *this,
        dispatch_key,
        std::move(kernel),
        // NOLINTNEXTLINE(performance-move-const-arg)
        std::move(cpp_signature),
        std::move(inferred_function_schema),
        std::move(debug)
    );

    ++op.operatorDef_->def_and_impl_count;

    return RegistrationHandlerRAII([this, op, op_name, dispatch_key, handle] {
        deregisterImpl_(op, op_name, dispatch_key, handle);
    });
}

```

```

// aten/src/ATen/core/dispatch/OperatorEntry.cpp
OperatorEntry::AnnotatedKernelContainerIterator OperatorEntry::registerKernel(
    const c10::Dispatcher& dispatcher,
    c10::optional<DispatchKey> dispatch_key,
    KernelFunction kernel,
    c10::optional<CppSignature> cpp_signature,
    std::unique_ptr<FunctionSchema> inferred_function_schema,
    std::string debug
) {

    //  cpp_signature

    //  schema

    // Add the kernel to the kernels list,
    // possibly creating the list if this is the first kernel.
    // Redirect catchAll registrations to CompositeImplicitAutograd.
    auto& k = dispatch_key.has_value() ? kernels_[*dispatch_key] : kernels_[DispatchKey::CompositeImplicitAutograd];

    //  dispatch key,

    //  kernel    OperatorEntry dispatch key
#ifdef C10_DISPATCHER_ONE_KERNEL_PER_DISPATCH_KEY
    k[0].kernel = std::move(kernel);
    k[0].inferred_function_schema = std::move(inferred_function_schema);
    k[0].debug = std::move(debug);
#else
    k.emplace_front(std::move(kernel), std::move(inferred_function_schema), std::move(debug));
#endif

    //  dispatch table
    AnnotatedKernelContainerIterator inserted = k.begin();
    // update the dispatch table, i.e. re-establish the invariant
    // that the dispatch table points to the newest kernel
    if (dispatch_key.has_value()) {
        updateDispatchTable_(dispatcher, *dispatch_key);
    } else {
        updateDispatchTableFull_(dispatcher);
    }
    return inserted;
}

```

算子封装

前面介绍到，注册算子的CPU实现的时候，注册的是函数wrapper_add_Tensor:

```
// build/aten/src/ATen/RegisterCPU.cpp
```

```
at::Tensor wrapper_add_Tensor(const at::Tensor & self, const at::Tensor & other, const at::Scalar & alpha) {
    structured_ufunc_add_CPU_functional op;
    op.meta(self, other, alpha);
    op.impl(self, other, alpha, *op.outputs_[0]);
    return std::move(op.outputs_[0]).take();
}
```

其中meta函数会调用到命名空间meta下的函数，其中TORCH_META_FUNC2(add, Tensor)等同于“void structured_add_Tensor::meta”。

```
// aten/src/ATen/native/BinaryOps.cpp
```

```
namespace meta {
```

```
TORCH_META_FUNC2(add, Tensor) (
    const Tensor& self, const Tensor& other, const Scalar& alpha
) {
    build_borrowing_binary_op(maybe_get_output(), self, other);
    native::alpha_check(dtype(), alpha);
}
```

在

```
// build/aten/src/ATen/UfuncCPUkernel_add.cpp
```

```
void add_kernel(TensorIteratorBase& iter, const at::Scalar & alpha) {
    at::ScalarType st = iter.common_dtype();
    RECORD_KERNEL_FUNCTION_DTYPE("add_stub", st);
    switch (st) {

        AT_PRIVATE_CASE_TYPE("add_stub", at::ScalarType::Bool, bool,
            [&]() {

                auto _s_alpha = alpha.to<scalar_t>();
                cpu_kernel(iter,
                    [=](scalar_t self, scalar_t other) { return ufunc::add(self, other, _s_alpha); }
                );
            }
        );
    }
}
```

算子注册:

```
// build/aten/src/ATen/UfuncCPUkernel_add.cpp
```

```
using add_fn = void (*)(TensorIteratorBase&, const at::Scalar &);
```

```
DECLARE_DISPATCH(add_fn, add_stub);
REGISTER_DISPATCH(add_stub, &add_kernel);
```

// aten/src/ATen/native/DispatchStub.cpp

```
#define DECLARE_DISPATCH(fn, name) \
    struct name : DispatchStub<fn, name> { \
        name() = default; \
        name(const name&) = delete; \
        name& operator=(const name&) = delete; \
    }; \
    extern TORCH_API struct name name

#define REGISTER_DISPATCH(name, fn) REGISTER_ARCH_DISPATCH(name, CPU_CAPABILITY, fn)
#define REGISTER_ARCH_DISPATCH(name, arch, fn) \
    template <> name::FnPtr TORCH_API DispatchStub<name::FnPtr, struct name>::arch = fn;
```

OperatorHandle

这里看到两种注册的类型，一种是OperatorHandler，注册到operatorLookupTable_中，可以根据OperatorName查询

比如对于例子中的 $y = x + 2$ 这条语句，dispatcher会查询到一个OperatorHandler

op，op.operatorDef_>op.name_就是OperatorName(“aten::add”，“Tensor”)，但是注册的kernelfunction很

// aten/src/ATen/core/dispatch/Dispatcher.h

```
class TORCH_API OperatorHandle {
public:
    OperatorHandle(OperatorHandle&&) noexcept = default;
    // ...

    // See [Note: Argument forwarding in the dispatcher] for why Args doesn't use C++
    C10_ALWAYS_INLINE Return call(Args... args) const {
        return c10::Dispatcher::singleton().call<Return, Args...>(*this, std::forward<Args>(args))
    }

    // ...

private:
    // ...
    Dispatcher::OperatorDef* operatorDef_;
    std::list<Dispatcher::OperatorDef>::iterator operatorIterator_;
};
```

OperatorHandle的call()方法会调用Dispatcher::call()方法。

继续跟踪，会走到

```
at::native::AVX2::cpu_kernel_vec<> (grain_size=32768, vop=..., op=..., iter=...)
  at ../aten/src/ATen/native/cpu/Loops.h:349
```

```
#0 at::native::AVX2::cpu_kernel_vec<> (grain_size=32768, vop=..., op=..., iter=...)
  at ../aten/src/ATen/native/cpu/Loops.h:349
#1 at::native::(anonymous namespace)::<lambda()>::operator() (_closure=<optimized out>)
  at /lab/tmp/pytorch/build/aten/src/ATen/UfuncCPUKernel_add.cpp:61
#2 at::native::(anonymous namespace)::add_kernel (iter=..., alpha=...)
  at /lab/tmp/pytorch/build/aten/src/ATen/UfuncCPUKernel_add.cpp:61
#3 0x00007fffe717e7be in at::(anonymous namespace)::wrapper_add_Tensor (self=..., other=...)
  at aten/src/ATen/RegisterCPU.cpp:1595
```

(gdb) bt

```
#0 at::native::AVX2::vectorized_loop<at::native::(anonymous namespace)::add_kernel(at::Tensor, at::Tensor, double, bool)> (at=0x7fffd300, at=0x7fffd300, alpha=1.0, inplace=false)
  at ../aten/src/ATen/native/cpu/Loops.h:212
#1 at::native::AVX2::VectorizedLoop2d<at::native::(anonymous namespace)::add_kernel(at::Tensor, at::Tensor, double, bool)> (at=0x7fffd300, at=0x7fffd300, alpha=1.0, inplace=false)
  at ../aten/src/ATen/native/cpu/Loops.h:287
#2 at::native::AVX2::unroll_contiguous_scalar_checks<function_traits<at::native::(anonymous namespace)::add_kernel(at::Tensor, at::Tensor, double, bool)>> (cb=..., strides=0x7fffd300) at ../aten/src/ATen/native/cpu/Loops.h:246
#3 at::native::AVX2::unroll_contiguous_scalar_checks<function_traits<at::native::(anonymous namespace)::add_kernel(at::Tensor, at::Tensor, double, bool)>> (cb=..., strides=0x7fffd300) at ../aten/src/ATen/native/cpu/Loops.h:248
#4 at::native::AVX2::VectorizedLoop2d<at::native::(anonymous namespace)::add_kernel(at::Tensor, at::Tensor, double, bool)> (at=0x7fffd300, at=0x7fffd300, alpha=1.0, inplace=false)
  at ../aten/src/ATen/native/cpu/Loops.h:283
#5 c10::function_ref<void(char**, long int const*, long int, long int)>::callback_fn<at::native::(anonymous namespace)::add_kernel(at::Tensor, at::Tensor, double, bool)> (params#0=params#0@entry=0x7fffd300, params#1=params#1@entry=0x7fffd300, params#2=params#2@entry=0x7fffd300, params#3=params#3@entry=1) at ../c10/util/FunctionRef.h:43
```

算子调用的过程

我们再看一个简单的例子：

```
import torch
```

```
x = torch.randn(2,2, requires_grad=True)
```

```
y = x + 2
```

在调用上，依次进行如下的调用：

```
// torch/csrc/autograd/generated/python_variable_methods.cpp
//      self_args      kwargs = 0x0
static PyObject * THPVariable_add(PyObject* self_, PyObject* args, PyObject* kwargs);

// torch/csrc/autograd/generated/python_arg_parser.h
```

```

inline PythonArgs PythonArgParser::parse(PyObject* self, PyObject* args, PyObject* kwargs)

// torch/csrc/utils/python_arg_parser.cpp
PythonArgs PythonArgParser::raw_parse(PyObject* self, PyObject* args, PyObject* kwargs)
bool FunctionSignature::parse(PyObject* self, PyObject* args, PyObject* kwargs, PyObject*

// torch/include/ATen/core/TensorBody.h --- generated from aten/src/ATen/templates/TensorBody.h
inline at::Tensor & Tensor::add(const at::Tensor & other, const at::Scalar & alpha) const

// build/aten/src/ATen/Operators_2.cpp
at::Tensor & add_Tensor::call(at::Tensor & self, const at::Tensor & other, const at::Scalar & alpha) const

// aten/src/ATen/core/dispatch/Dispatcher.cpp
OperatorHandle Dispatcher::findSchemaOrThrow(const char* name, const char* overloadName) const
c10::optional<OperatorHandle> Dispatcher::findSchema(const OperatorName& name, const char* overloadName) const
c10::optional<OperatorHandle> Dispatcher::findOp(const OperatorName& name, const char* overloadName) const

// aten/src/ATen/core/dispatch/Dispatcher.cpp
Return TypedOperatorHandle::call(Args... args) const;

// aten/src/ATen/core/dispatch/Dispatcher.cpp
Return Dispatcher::call(const TypedOperatorHandle<Return(Args...)>& op, Args... args) const

// aten/src/ATen/core/dispatch/DispatchKeyExtractor.h
DispatchKeySet DispatchKeyExtractor::getDispatchKeySetUnboxed(const Args... args) const

// aten/src/ATen/core/boxing/KernelFunction.h
Return call(const OperatorHandle& opHandle, DispatchKeySet dispatchKeySet, Args... args) const

// torch/csrc/autograd/utils/wrap_outputs.h
// Python
PyObject* wrap(PyTypeObject *type, std::tuple<Ts...> values);

```

在进入C++层面的第一步，是进行调用参数的解码。因为在Python层面和在C++层面类的体系是不一样的，Python语言PyTorch为此定义了PythonArgParser类，在函数被调用的入口处进行参数解析：

```

// torch/csrc/autograd/generated/python_variable_methods.cpp
static PyObject * THPVariable_add(PyObject* self_, PyObject* args, PyObject* kwargs)
{
    HANDLE_TH_ERRORS
    const Tensor& self = THPVariable_Unpack(self_);
    static PythonArgParser parser({
        "add(Scalar alpha, Tensor other)|deprecated",
        "add(Tensor other, *, Scalar alpha=1)",
    }, /*traceable=*/true);
}

```

```

ParsedArgs<2> parsed_args;
auto _r = parser.parse(self_, args, kwargs, parsed_args);
if(_r.has_torch_function()) {
    return handle_torch_function(_r, self_, args, kwargs, THPVariableClass, "torch.Tensor");
}
switch (_r.idx) {
case 0: {
    // [deprecated] aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor

    auto dispatch_add = [](const at::Tensor & self, const at::Scalar & alpha, const at::Tensor & other) {
        pybind11::gil_scoped_release no_gil;
        return self.add(other, alpha);
    };
    return wrap(dispatch_add(self, _r.scalar(0), _r.tensor(1)));
}
case 1: {
    // aten::add.Tensor(Tensor self, Tensor other, *, Scalar alpha=1) -> Tensor

    auto dispatch_add = [](const at::Tensor & self, const at::Tensor & other, const at::Scalar & alpha) {
        pybind11::gil_scoped_release no_gil;
        return self.add(other, alpha);
    };
    return wrap(dispatch_add(self, _r.tensor(0), _r.scalar(1)));
}
}
Py_RETURN_NONE;
END_HANDLE_TH_ERRORS
}

```

如上面的代码，对于add方法，Pytorch支持两种不同的签名，但是前一种已经过时了，因此实际调用走的都是第二种。在C API: PyTuple_GET_ITEM()和PyDict_GetItem()，在调用Tensor::add()之前，PythonArgParser会通过其tensor()和Tensor::add()的实现中，并不是真正的算子代码，因为刚才只完成了从Python到C++的调用转换，实际的算子实现

```

// torch/include/ATen/core/TensorBody.h

```

```

// aten::add_.Tensor(Tensor(a!) self, Tensor other, *, Scalar alpha=1) -> Tensor(a!)
inline at::Tensor & Tensor::add_(const at::Tensor & other, const at::Scalar & alpha) const {
    return at::_ops::add__Tensor::call(const_cast<Tensor&>(*this), other, alpha);
}

```

```

// build/aten/src/ATen/Operators_2.cpp

```

```

STATIC_CONST_STR_OUT_OF_LINE_FOR_WIN_CUDA(add__Tensor, name, "aten::add_")
STATIC_CONST_STR_OUT_OF_LINE_FOR_WIN_CUDA(add__Tensor, overload_name, "Tensor")

```

```

STATIC_CONST_STR_OUT_OF_LINE_FOR_WIN_CUDA(add__Tensor, schema_str, "add_.Tensor(Tensor(a!) s

// aten::add_.Tensor(Tensor(a!) self, Tensor other, *, Scalar alpha=1) -> Tensor(a!)
static C10_NOINLINE c10::TypedOperatorHandle<add__Tensor::schema> create_add__Tensor_typed_h
    return c10::Dispatcher::singleton()
        .findSchemaOrThrow(add__Tensor::name, add__Tensor::overload_name)
        .typed<add__Tensor::schema>();
}

// aten::add_.Tensor(Tensor(a!) self, Tensor other, *, Scalar alpha=1) -> Tensor(a!)
at::Tensor & add__Tensor::call(at::Tensor & self, const at::Tensor & other, const at::Scalar

    static auto op = create_add__Tensor_typed_handle();
    return op.call(self, other, alpha);
}

THPVariable_add ->

```

自定义算子的实现过程

原生算子的实现

所谓“原生”，指的就是内置在PyTorch中的算子，跟随PyTorch一起编译生成，可以同“torch.xxx”等方式使用的。由于原生算子的数量非常多，处于效率和可用性的考虑，在不同的平台上可能会有实现，另外算子要支持注册到torch.nn。很多原生算子的模板定义在native_functions.yaml中，比如sigmoid函数：

```

# aten/src/ATen/native/native_functions.yaml

- func: sigmoid(Tensor self) -> Tensor
  device_check: NoCheck # TensorIterator
  structured_delegate: sigmoid.out
  variants: function, method
  dispatch:
    QuantizedCPU: sigmoid_quantized_cpu
    MklDnnCPU: mklDnn_sigmoid

- func: sigmoid_backward(Tensor grad_output, Tensor output) -> Tensor
  python_module: nn
  structured_delegate: sigmoid_backward.grad_input

```

其中： - func字段定义了算子的名称和输入输出参数。 - device_check: 暂时还不清楚用途，在模板里都是NoCheck。 - structured_delegate: sigmoid.out - variants字段生命这个算子的类型和使用方式，function表明sigmoid这个算子可以通过函数torch.sigm

- dispatch字段定义了在不同的平台或者优化方式下该算子的变体。这里针对使用量化方式运行时，会调用相应的量
- python-module字段定义了该算法会被注册到的Python模块。

sigmoid函数是机器学习中最基本的函数之一，其公式如下：

$$\text{Sigmoid}(x) = \frac{1}{1 + e^{-x}}$$

我们在使用sigmoid函数时，调用的是torch.nn.Sigmoid函数，其背后则是调用了torch.sigmoid()函数，也就是上面

```
class Sigmoid(Module):
    r"""Applies the element-wise function:
    Examples::
        >>> m = nn.Sigmoid()
        >>> input = torch.randn(2)
        >>> output = m(input)
    """
```

```
def forward(self, input: Tensor) -> Tensor:
    return torch.sigmoid(input)
```

在tools/autograd/derivatives.yaml中，定义了算子的前向计算输出反向计算梯度的对应关系，比如sigmoid算子的

```
- name: sigmoid(Tensor self) -> Tensor
  self: sigmoid_backward(grad, result)
  result: auto_element_wise
```

在native_functions.yaml中只是声明了sigmoid算子，具体的算子实现是和平台相关的，因此要到各个平台目录下去

```
// aten/src/ATen/native/cpu/UnaryOpsKernel.cpp
```

```
static void sigmoid_kernel(TensorIteratorBase& iter) {
    if (iter.common_dtype() == kBFloat16) {
        cpu_kernel_vec(
            iter,
            [=](BFloat16 a) -> BFloat16 {
                float a0 = static_cast<float>(a);
                return static_cast<float>(1) / (static_cast<float>(1) + std::exp((-a0)));
            },
            [=](Vectorized<BFloat16> a) {
                Vectorized<float> a0, a1;
                std::tie(a0, a1) = convert_bfloat16_float(a);
                a0 = (Vectorized<float>(static_cast<float>(1)) + a0.neg().exp()).reciprocal();
                a1 = (Vectorized<float>(static_cast<float>(1)) + a1.neg().exp()).reciprocal();
                return convert_float_bfloat16(a0, a1);
            });
    } else {
        AT_DISPATCH_FLOATING_AND_COMPLEX_TYPES(iter.common_dtype(), "sigmoid_cpu", [&]() {
            cpu_kernel_vec(
```

```

        iter,
        [=](scalar_t a) -> scalar_t {
            return (static_cast<scalar_t>(1) / (static_cast<scalar_t>(1) + std::exp((-a))));
        },
        [=](Vectorized<scalar_t> a) {
            a = Vectorized<scalar_t>(static_cast<scalar_t>(0)) - a;
            a = a.exp();
            a = Vectorized<scalar_t>(static_cast<scalar_t>(1)) + a;
            a = a.reciprocal();
            return a;
        });
    });
}
}

REGISTER_DISPATCH(sigmoid_stub, &CPU_CAPABILITY::sigmoid_kernel);

// aten/src/ATen/native/cpu/BinaryOpsKernel.cpp

void sigmoid_backward_kernel(TensorIteratorBase& iter) {
    if (isComplexType(iter.dtype())) {
        // .....
    } else if (iter.dtype() == kBFloat16) {
        // .....
    } else {
        // .....
    }
}

// aten/src/ATen/native/cpu/UnaryOps.cpp

CREATE_UNARY_FLOAT_META_FUNC(sigmoid)

CREATE_UNARY_TORCH_IMPL_FUNC(sigmoid_out, sigmoid_stub)
DEFINE_DISPATCH(sigmoid_stub); // NOLINT(cppcoreguidelines-avoid-non-const-global-variables)

```

在sigmoid_kernel()的实现里，根据传输Tensor类型的不同，构建了不同的匿名函数，然后调用cpu_kernel_vec()。sigmoid_kernel是sigmoid算子在cpu下的实现，当然即使在CPU下，sigmoid函数也有多种形式，除了普通的浮点计算。AT_DISPATCH_FLOATING_AND_COMPLEX_TYPES宏有三个参数：- iter.common_dtype()，指明操作的Tensor属于哪种类型 - “sigmoid_cpu”，算子的名称 - 匿名函数，调用了cpu_kernel_vec

在aten/src/ATen/native/cpu/Loops.cpp中，有两个cpu_kernel相关的函数，由于cpu下的源文件在编译的时候会加上-cpu_kernel()：依赖于编译器自动实现计算的向量化 - cpu_kernel_vec()：使用x86 SIMD原语实现向量化。一般来讲，使用cpu_kernel_vec()的时候，说明实现该算子的实现是经过精心优化的，效率

例如用这两个函数实现浮点数相乘的算子，可以这样实现：

```
cpu_kernel(iter, [](float a, float b) { return a * b; });

cpu_kernel_vec(iter,
    [](float a, float b) { return a * b; },
    [](Vectorized<float> a, Vectorized<float> b) { return a * b; });
```

下面我们看一下cpu_kernel_vec()函数的实现：

```
// aten/src/ATen/native/cpu/Loops.cpp
```

```
template <bool check_dynamic_cast=true, typename func_t, typename vec_func_t>
void cpu_kernel_vec(TensorIteratorBase& iter, func_t&& op, vec_func_t&& vop, int64_t grain_size) {
    using traits = function_traits<func_t>;
    // this could be extended to work with void return types
    TORCH_INTERNAL_ASSERT(iter.ninputs() == traits::arity);
    TORCH_INTERNAL_ASSERT(iter.noutputs() == 1);
    // dynamic casting not currently supported on CPU, but some kernels (like Fill)
    // explicitly dynamic_cast, so we give the opt-out of checking.
    c10::guts::if_constexpr<check_dynamic_cast>([&] {
        TORCH_INTERNAL_ASSERT(!needs_dynamic_casting<func_t>::check(iter));
    });

    iter.for_each(make_vectorized_loop2d(op, vop), grain_size);
    iter.cast_outputs();
}
```

可以看到，对每个Tensor，又调用了make_vectorized_loop2d()

```
// aten/src/ATen/native/cpu/Loops.cpp
```

```
template <typename op_t, typename vop_t>
VectorizedLoop2d<op_t, vop_t> make_vectorized_loop2d(
    const op_t &op, const vop_t &vop) {
    return VectorizedLoop2d<op_t, vop_t>(op, vop);
}
```

```
template <typename op_t, typename vop_t>
struct VectorizedLoop2d {
    op_t op;
    vop_t vop;

    using traits = function_traits<op_t>;
    static constexpr int ntensors = traits::arity + 1;
    using data_t = std::array<char*, ntensors>;

    VectorizedLoop2d(const op_t &op, const vop_t &vop):
```

```

    op(op), vop(vop) {}

static void advance(data_t &data, const int64_t *outer_strides) {
    for (const auto arg : c10::irange(data.size())) {
        data[arg] += outer_strides[arg];
    }
}

void operator()(char** base, const int64_t *strides, int64_t size0, int64_t size1) {
    data_t data;
    std::copy_n(base, ntensors, data.data());
    const int64_t *outer_strides = &strides[ntensors];

    if (is_contiguous<traits>(strides)) {
        for (const auto i : c10::irange(size1)) {
            (void)i;
            vectorized_loop(data.data(), size0, 0, op, vop);
            advance(data, outer_strides);
        }
    } else {
        using Indices = std::make_index_sequence<traits::arity>;
        unroll_contiguous_scalar_checks<traits>(strides, Indices{}, [&](size_t idx) {
            if (idx) {
                for (const auto i : c10::irange(size1)) {
                    (void)i;
                    vectorized_loop(data.data(), size0, idx, op, vop);
                    advance(data, outer_strides);
                }
            } else {
                for (const auto i : c10::irange(size1)) {
                    (void)i;
                    basic_loop(data.data(), strides, 0, size0, op);
                    advance(data, outer_strides);
                }
            }
        });
    }
}
};

```

很明显，VectorizedLoop2d的主要工作就是根据Tensor的stride的不同，选择不同的调用模式，但最终不管是调用v现在我们回到当初sigmoid函数的实现部分，其中对每个Tensor的操作函数实现是这样的：

```

// aten/src/ATen/native/cpu/UnaryOpsKernel.cpp
cpu_kernel_vec(
    iter,

```

```

[=](scalar_t a) -> scalar_t {
    return (static_cast<scalar_t>(1) / (static_cast<scalar_t>(1) + std::exp((-a))));
},
[=](Vectorized<scalar_t> a) {
    a = Vectorized<scalar_t>(static_cast<scalar_t>(0)) - a;
    a = a.exp();
    a = Vectorized<scalar_t>(static_cast<scalar_t>(1)) + a;
    a = a.reciprocal();
    return a;
});

// aten/src/ATen/native/cpu/vec/vec256/vec256_float.h
Vectorized<float> exp() const {
    return Vectorized<float>(Sleef_expf8_u10(values));
}

```

在代码中可以看出，对应cpu的实现有很多，实际运行时会根据不同的平台和数据类型调用相应的实现，以达到比较

<https://blog.csdn.net/yelede2009/article/details/120411361>

有各种函数库以向量方式来计算数学函数，例如：对数、幂函数、三角函数等。这些函数库对向量化数学代码有两种不同种类的向量数学库：长向量库和短向量库。来看看它们的不同。假设要计算1000个数字的某个函数，一个库函数存储这1000个结果到另一个数组。使用长向量版库函数的缺点是，如果要做一系列计算，在下一次调用的向量库，可以把数据集拆分为子向量来适配向量寄存器。如果向量寄存器可以处理4个数字，那么需要调用250次，被下一次计算利用，而不需要存储中间结果到RAM中。这可能更快。然而，短向量的库函数可能是不利的，如果要做一系列计算，这是一些长向量函数库：

Intel 向量数学库（VML, MKL）。工作在x86平台。这些库函数在非Intel的CPU上会低效，除非重写了Intel cpu分发器。Intel的IPP。工作在x86平台。也适用于非Intel的CPU。包含很多统计、信号处理和图像处理函数。Yeppp。开源库。支持x86和ARM平台，多种编程语言。参考Yeppp。

这是一些短向量库：

Sleef库。支持多种平台。开源。参考www.sleef.org。Intel短向量库（SVML）。Intel编译器提供，被自动向量化。mveclibabi=svml使用这个库。如果用的是非Intel的CPU，也可以使用。

AMD LIBM库。只支持64位Linux平台。没有FMA4指令集时，性能会降低。Gnu通过-mveclibabi=acml选项使用。VCL库。个人开发。参考<https://github.com/vectorclass>。

Dispatch的过程似乎有些复杂，有很多宏处理，更是导致不容易看懂。 “ ‘C++ //
aten/src/ATen/Dispatch.h

```

define AT_PRIVATE_CASE_TYPE(NAME, enum_type, type,
...)
AT_PRIVATE_CASE_TYPE_USING_HINT(NAME, enum_type,
type, scalar_t, VA_ARGS)

define AT_DISPATCH_FLOATING_TYPES_AND_HALF(TYPE,
NAME, ...)
[&] {
const auto& the_type = TYPE;
/* don't use TYPE again in case it is an expensive
or side-effect op */
at::ScalarType _st = ::detail::scalar_type(the_type);
RECORD_KERNEL_FUNCTION_DTYPE(NAME, _st);
switch (_st) {
AT_PRIVATE_CASE_TYPE(NAME, at::ScalarType::Double,
double, VA_ARGS)
AT_PRIVATE_CASE_TYPE(NAME, at::ScalarType::Float,
float, VA_ARGS)
AT_PRIVATE_CASE_TYPE(NAME, at::ScalarType::Half,
at::Half, VA_ARGS)
default:
AT_ERROR(#NAME, " not implemented for " ,
toString(_st), " " );
}
}() " "

```

宏AT_DISPATCH_FLOATING_AND_COMPLEX_TYPES

参考

- <https://pytorch.org/tutorials/advanced/dispatcher.html>
- <http://blog.ezyang.com/2020/09/lets-talk-about-the-pytorch-dispatcher/>

- https://blog.csdn.net/Chris_zhangrx/article/details/119512418
- <https://zhuanlan.zhihu.com/p/67834038>
- <https://blog.csdn.net/xixiaoyaoww/article/details/112211025>
- pytorch中的dispatcher <https://zhuanlan.zhihu.com/p/390049109>
- [Pytorch 源码阅读] —— 谈谈 dispatcher (二) https://blog.csdn.net/Chris_zhangrx/article/details/
- [Pytorch 源码阅读] —— 谈谈 dispatcher (一) https://blog.csdn.net/Chris_zhangrx/article/details/
- <https://zhuanlan.zhihu.com/p/349560723>
- <https://zhuanlan.zhihu.com/p/499979372>
- 这可能是关于Pytorch底层算子扩展最详细的总结了 <https://wenku.baidu.com/view/1415b43ac181e53a58021>

计算图

基本内容

本章内容主要回答以下几个问题：

神经网络的基本结构

深度学习框架时如何执行计算图的

计算图执行过程中的基本数据结构

PyTorch中的具体实现

神经网络的基本结构

深度学习解决的是深度神经网络的优化问题，虽然深度神经网络的模型种类繁多，从最简单的MLP模型到近年流行的

```
import torch
from torch import nn

class DemoNet(nn.Module):
    def __init__(self):
        super(DemoNet, self).__init__()
        self.w = torch.rand(2,2)
    def forward(self, x):
        y = self.w * x
        return y * y

input = torch.rand(2, 2)
model = DemoNet()
```

使用TensorBoard查看该网络的可视化，如下图：

其中y处是一个算子” Operation: aten::mul “

虽然上面只是最简单的一个例子，但也包括了神经网络作为有向无环图的基本结构：

- 顶点：代表一个输入数据、算子、或者输出数据 - 边：代表数据和算子、算子和算子之间的输入输出关系。

深度神经网络包括结果的前向计算过程和梯度的反向传播过程，显而易见的是，深度学习框架需要事先构造计算图，

- 根据代码逻辑，构造好一个计算图，之后这个计算图可以反复执行 - 每次在执行时，都重新构造好计算图

PyTorch选择的是第二种方式，也就是动态图的方式。动态图的好处是可以在代码逻辑中使用各种条件判断。

PyTorch中计算图的实现

虽然不是所有的计算图都通过上面的例子中的nn.Module来实现，但nn.Module确实是PyTorch中神经网络的基础结构

```
# torch/nn/modules/module.py
```

```
class Module:
    """Base class for all neural network modules.
    ...
    """

    training: bool
    _is_full_backward_hook: Optional[bool]

    def __init__(self) -> None:
        """
        Initializes internal Module state, shared by both nn.Module and ScriptModule.
        """

        torch._C._log_api_usage_once("python.nn.module")

        self.training = True
        self._parameters: Dict[str, Optional[Parameter]] = OrderedDict()
        self._buffers: Dict[str, Optional[Tensor]] = OrderedDict()
        self._non_persistent_buffers_set: Set[str] = set()
        self._backward_hooks: Dict[int, Callable] = OrderedDict()
        self._is_full_backward_hook = None
        self._forward_hooks: Dict[int, Callable] = OrderedDict()
        self._forward_pre_hooks: Dict[int, Callable] = OrderedDict()
        self._state_dict_hooks: Dict[int, Callable] = OrderedDict()
        self._load_state_dict_pre_hooks: Dict[int, Callable] = OrderedDict()
        self._load_state_dict_post_hooks: Dict[int, Callable] = OrderedDict()
        self._modules: Dict[str, Optional['Module']] = OrderedDict()

        forward: Callable[..., Any] = _forward_unimplemented
```

Module类的主要属性及方法如下：

一个神经网络，最重要的是其内部的参数，在Module中有两个属性和参数相关：_parameters和_buffers，它们的类从定义上看，_buffers中存放的是Tensor类型的数据，而_parameters中存放的是Parameter类型的数据，在构造时参

```
# torch/nn/parameter.py
```

```
class Parameter(torch.Tensor, metaclass=_ParameterMeta):
    def __new__(cls, data=None, requires_grad=True):
        # .....
```

当构造好Parameter并且赋值给nn.Module时，会自动调用nn.Module的register_parameter()方法进行注册。

```
# torch/nn/modules/module.py
```

```
class Module:

    def __setattr__(self, name: str, value: Union[Tensor, 'Module']) -> None:

        params = self.__dict__.get('_parameters')
        if isinstance(value, Parameter):
            self.register_parameter(name, value)
        # handle value with other types
```

为了看的更清楚一些，我们看一下PyTorch中内置的网络组件，例如：

```
# torch/nn/modules/conv.py
```

```
class _ConvNd(Module):

    __constants__ = ['stride', 'padding', 'dilation', 'groups',
                    'padding_mode', 'output_padding', 'in_channels',
                    'out_channels', 'kernel_size']
    __annotations__ = {'bias': Optional[torch.Tensor]}

    def _conv_forward(self, input: Tensor, weight: Tensor, bias: Optional[Tensor]) -> Tensor
        ...

    _in_channels: int
    _reversed_padding_repeated_twice: List[int]
    out_channels: int
    kernel_size: Tuple[int, ...]
    stride: Tuple[int, ...]
    padding: Union[str, Tuple[int, ...]]
    dilation: Tuple[int, ...]
    transposed: bool
    output_padding: Tuple[int, ...]
    groups: int
    padding_mode: str
    weight: Tensor
```

```

bias: Optional[Tensor]

def __init__(self,
              in_channels: int,
              out_channels: int,
              kernel_size: Tuple[int, ...],
              stride: Tuple[int, ...],
              padding: Tuple[int, ...],
              dilation: Tuple[int, ...],
              transposed: bool,
              output_padding: Tuple[int, ...],
              groups: int,
              bias: bool,
              padding_mode: str,
              device=None,
              dtype=None) -> None:
    super(_ConvNd, self).__init__()

    # check and handle padding and other parameter...

    if transposed:
        self.weight = Parameter(torch.empty(
            (in_channels, out_channels // groups, *kernel_size), **factory_kwargs))
    else:
        self.weight = Parameter(torch.empty(
            (out_channels, in_channels // groups, *kernel_size), **factory_kwargs))
    if bias:
        self.bias = Parameter(torch.empty(out_channels, **factory_kwargs))
    else:
        self.register_parameter('bias', None)

    self.reset_parameters()

```

计算图的执行过程

在深度学习中，我们的神经网络一般是基于nn.Module实现的，典型的调用方式是：

```

y = DemoNet(x)
loss = compute_loss(y, label)

```

可见计算图的执行其实就是nn.Module的调用过程，从下面的实现中可以看出，主要的工作就是调用forward()方法

```

# torch/nn/modules/module.py

```

```

class Module:

```



```

def _call_impl(self, *input, **kwargs):
    forward_call = (self._slow_forward if torch._C._get_tracing_state() else self.forward)

    # YL: handle pre-forward hooks, you can change input here
    # ...

    result = forward_call(*input, **kwargs)
    # YL: handle forward hooks
    # ...

    # Handle the non-full backward hooks
    # ...

    return result

```

```
__call__ : Callable[..., Any] = _call_impl
```

相应的，我们可以看一下卷积操作的实现：

```
# torch/nn/modules/conv.py
```

```
from .. import functional as F
```

```
class Conv2d(_ConvNd):
```

```
    ## YL __init__() implemetation here
```

```

def _conv_forward(self, input: Tensor, weight: Tensor, bias: Optional[Tensor]):
    if self.padding_mode != 'zeros':
        return F.conv2d(F.pad(input, self._reversed_padding_repeated_twice, mode=self.padding_mode),
                        weight, bias, self.stride,
                        _pair(0), self.dilation, self.groups)
    return F.conv2d(input, weight, bias, self.stride,
                    self.padding, self.dilation, self.groups)

```

```

def forward(self, input: Tensor) -> Tensor:
    return self._conv_forward(input, self.weight, self.bias)

```

由此可见，卷积算子的实现调用了functional模块中的卷积函数。这也说明，在PyTorch中，神经网络的定义和算子Dispatch

参考

- <https://zhuanlan.zhihu.com/p/89442276>

自动微分

自动微分一直被视为深度学习框架的核心能力，在训练深度学习神经网络的时候，网络的参数需要根据输出端的梯度

自动微分的理论基础

在了解自动微分之前，我们先从优化的角度看一下参数和梯度的关系，这也是深度学习的目标。

考虑下面这个公式，这是典型的线性回归的公式，我们需要根据输出与实际值的差异调整系数 w 及截距 b ：

$$y = w * x + b$$

根据微分原理我们知道：

$$\frac{\partial y}{\partial w} = x$$
$$\frac{\partial y}{\partial b} = 1$$

根据上面的式子，在微小的取值范围内，为了调整 w ，可以这样计算：

$$dw = x * dy$$

其中 dy 就是输出与实际值的差异。在实际计算中，由于 dy 的值不会很小，我们会加一个比较小的系数 α 来缓慢调整

$$dw = \alpha * x * dy$$

同理，对于另一个算子：

$$y = w * x^2$$

我们可以计算得到：

$$dw = \alpha * x^2 * dy$$

下面我们看看自动微分是怎样在PyTorch中实现的，在探究之前，我们先关注几个问题：

- PyTorch中的计算图是怎样构建的？
- 反向传播的流程是什么样的？

计算图及反向传播

在计算图中，autograd会记录所有的操作，并生成一个DAG（有向无环图），其中输出的tensor是根节点，输入的te

在前向阶段，autograd同时做两件事： - 根据算子计算结果Tensor - 维护算子的梯度函数

在反向阶段，当.backward()被调用时，autograd： - 对于节点的每一个梯度函数，计算相应节点的梯度

- 在节点上对梯度进行累加，并保存到节点的.grad属性上
- 根据链式法则，按照同样的方式计算，一直到叶子节点

对于一个简单的例子：

```
import torch

a = torch.tensor([2., 3.], requires_grad=True)
b = torch.tensor([6., 4.], requires_grad=True)

Q = 3*a**3 - b**2
```

下图是对应的计算图，其中的函数代表梯度计算函数：

自动微分相关的核心数据结构

TensorImpl是Tensor的实现

at::Tensor: shared ptr 指向 TensorImpl

TensorImpl: 对 at::Tensor 的实现

```
[AutogradMetaInterface](c10::AutogradMetaInterface) autograd_meta_ tensor variable
Variable: 就是Tensor，为了向前兼容保留的
using Variable = at::Tensor;

, Variable gradient, Tensor gradient

Variable AutogradMeta [AutogradMetaInterface](c10::AutogradMetaInterface) Variable
version view

AutogradMeta, autograd
// c10/core/TensorImpl.h

struct C10_API TensorImpl : public c10::intrusive_ptr_target {
    // ...
public:
    Storage storage_;

private:
    std::unique_ptr<c10::AutogradMetaInterface> autograd_meta_ = nullptr;

protected:
    std::unique_ptr<c10::NamedTensorMetaInterface> named_tensor_meta_ = nullptr;

    c10::VariableVersion version_counter_;

    PyObject* pyobj_;
```

```

c10::impl::SizesAndStrides sizes_and_strides_;

int64_t storage_offset_ = 0;

int64_t numel_ = 1;

caffe2::TypeMeta data_type_;

c10::optional<c10::Device> device_opt_;

bool is_contiguous_ : 1;

bool storage_access_should_throw_ : 1;

bool is_channels_last_ : 1;

bool is_channels_last_contiguous_ : 1;

bool is_channels_last_3d_ : 1;

bool is_channels_last_3d_contiguous_ : 1;

bool is_non_overlapping_and_dense_ : 1;

bool is_wrapped_number_ : 1;

bool allow_tensor_metadata_change_ : 1;

bool reserved_ : 1;

uint8_t sizes_strides_policy_ : 2;

DispatchKeySet key_set_;
}

```

autograd_meta_表示 Variable 中关于计算梯度的元数据信息，AutogradMetaInterface 是一个接口，有不同的子类，这里的 Variable 对象的梯度计算的元数据类型为 AutogradMeta，其部分成员为

```
// torch/csrc/autograd/variable.h
```

```

struct TORCH_API AutogradMeta : public c10::AutogradMetaInterface {
    std::string name_;

    Variable grad_;
    std::shared_ptr<Node> grad_fn_;
}

```

```

std::weak_ptr<Node> grad_accumulator_;
std::shared_ptr<ForwardGrad> fw_grad_;

std::vector<std::shared_ptr<FunctionPreHook>> hooks_;
std::shared_ptr<hooks_list> cpp_hooks_list_;

bool requires_grad_;
bool retains_grad_;
bool is_view_;
uint32_t output_nr_;

// ...
}

```

grad_ 表示反向传播时，关于当前 Variable 的梯度值。grad_fn_ 是用于计算非叶子-Variable 的梯度的函数，比如 AddBackward0 对象用于计算 result 这个 Variable 的梯度。对于叶子 Variable，此字段为 None。grad_accumulator_ 用于累加叶子 Variable 的梯度累加器，比如 AccumulateGrad 对象用于累加 self 的梯度。对于非叶子 Variable，此字段为 None。output_nr_ 表示当前 Variable 是 计算操作的第一个输出，此值从 0 开始。

可以看到，grad_fn_ 和 grad_accumulator_ 都是 Node 的指针，这是因为在计算图中，算子的 C++ 类型是 Node，不同的算子 Node 是由上一级的 Node 创建的

```

// torch/include/torch/csrc/autograd/function.h

struct TORCH_API Node : std::enable_shared_from_this<Node> {
public:
    /// Construct a new `Node` with the given `next_edges`
    // NOLINTNEXTLINE(cppcoreguidelines-pro-type-member-init)
    explicit Node(
        uint64_t sequence_nr,
        edge_list&& next_edges = edge_list()
        : sequence_nr_(sequence_nr),
        next_edges_(std::move(next_edges)) {

        for (const Edge& edge: next_edges_) {
            update_topological_nr(edge);
        }

        if (AnomalyMode::is_enabled()) {
            metadata()->store_stack();

            assign_parent();
        }

        // Store the thread_id of the forward operator.

```

```

    // See NOTE [ Sequence Numbers ]
    thread_id_ = at::RecordFunction::currentThreadId();
}

/// Evaluates the function on the given inputs and returns the result of the
/// function call.
variable_list operator()(variable_list&& inputs) {
    // ...
    return apply(std::move(inputs));
}

uint32_t add_input_metadata(const at::Tensor& t) noexcept {
    // ...
}

void add_next_edge(Edge edge) {
    update_topological_nr(edge);
    next_edges_.push_back(std::move(edge));
}

protected:
    /// Performs the `Node`'s actual operation.
    virtual variable_list apply(variable_list&& inputs) = 0;

    variable_list traced_apply(variable_list inputs);

    const uint64_t sequence_nr_;

    uint64_t topological_nr_ = 0;

    mutable bool has_parent_ = false;

    uint64_t thread_id_ = 0;

    std::mutex mutex_;

    edge_list next_edges_;

    PyObject* pyobj_ = nullptr;

```

```

std::unique_ptr<AnomalyMetadata> anomaly_metadata_ = nullptr;

std::vector<std::unique_ptr<FunctionPreHook>> pre_hooks_;

std::vector<std::unique_ptr<FunctionPostHook>> post_hooks_;

at::SmallVector<InputMetadata, 2> input_metadata_;
};

```

AutoGradMeta

```

AutoGradMeta : Variable autograd

```

```

grad_ Variable AutoGradMeta var tensor

```

```

Node grad_fn var graph grad_accumulator var , grad_

```

```

output_nr var grad_fn

```

```

Edge gradient_edge, gradient_edge.function grad_fn, gradient_edge.input_nr gr

```

Edge

autograd::Edge: 指向autograd::Node的一个输入

```

Node edge Node

```

```

input_nr edge Node

```

Node

autograd::Node: 对应AutoGrad Graph中的Op

```

autograd op apply

```

```

next_edges_

```

```

input_metadata_ tensor metadata

```

```

op

```

Node in AutoGrad Graph

```

Variable Edge Node

```

```

        Edge Var

call operator

next_edge

        Node

        Node next_edge(index)/next_edges()

        add_next_edge()

```

前向计算

PyTorch通过tracing只生成了后向AutoGrad Graph.

代码是生成的，需要编译才能看到对应的生成结果

```

gen_variable_type.py op

pytorch/torch/csrc/autograd/generated/

tracing

relu pytorch/torch/csrc/autograd/generated/VariableType_0.cpp

grad_fn trace op .

```

后向计算

```

autograd::backward():计算output var的梯度值，调用的 run_backward()

autograd::grad(): : 计算有output var和到特定input的梯度值，调用的
run_backward()

autograd::run_backward() • g' f

output var grad_fn roots

input var grad_fn output_edges,

autograd::Engine::get_default_engine().execute(...)

autograd::Engine::execute(...)

```



```

GraphTask

GraphRoot Node roots Node apply() roots grad

compute_dependencies(...)

GraphRoot grad_fn grad_fn GraphTask

GraphTask input var

GraphTask

CPU or GPU

CPU autograd::Engine::thread_main(...)
autograd::Engine::thread_main(...)
evaluate_function(...)

call_function(...) , Node

grad Tensor grad tensor grad_fn grad_fn backward backward

Topic .

```

参考

- <https://blog.csdn.net/zandaoguang/article/details/115713552>
- <https://zhuanlan.zhihu.com/p/111239415>
- <https://zhuanlan.zhihu.com/p/138203371>

数据加载

主要内容

数据的加载主要包括以下几个方面：

- 数据集的格式转换，需要支持各种类型各种格式的数据，如图片、语音、文本
- 数据的采样和shuffle，可能面临分布式的挑战。
- 数据增强，会产生额外的数据
- 数据预处理，如图片事先进行黑白二值化等
- 数据分batch
- 数据加载到内存，并且进入锁页内存
- 数据加载到GPU
- 数据分发给不同的计算单元，并且不会重复，且支持分布式训练

数据加载的设计

下面我们先看一个利用CIFAR10数据集进行模型训练的例子：

```
transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

trainset = torchvision.datasets.CIFAR10(
    root='./data', train=True, download=True, transform=transform_train)
trainloader = torch.utils.data.DataLoader(
    trainset, batch_size=128, shuffle=True, num_workers=2)

# Model
print('==> Building model..')
net = SENet18()
net = net.to(device)
if device == 'cuda':
    net = torch.nn.DataParallel(net)
    cudnn.benchmark = True

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=args.lr,
                       momentum=0.9, weight_decay=5e-4)
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=200)

# Training
def train(epoch):
    print('\nEpoch: %d' % epoch)
    net.train()
    train_loss = 0
    correct = 0
    total = 0
    for batch_idx, (inputs, targets) in enumerate(trainloader):
        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = net(inputs)
```

```
for epoch in range(start_epoch, start_epoch+200):
    train(epoch)
    scheduler.step()
```

并行数据读取

支持锁页内存

在深度学习模型训练过程中，因为数据集所占的内存比较多，又需要被频繁访问，因此一个比较好的加速方法就是利用锁页内存。使用锁页内存的另一个好处是主机内存和GPU内存之间的数据传输，基于锁页内存传输数据可以避免一次临时的数据

相比算子实现来讲，数据加载可以算作是非常简单直接的实现了。如下是单进程下数据加载的运行时，_SingleProc在多进程的情况下，最耗费时间的Fetcher部分和pin_memory()部分改成了多进程，如下图：

在torch模块中，DataSet是所有数据集的基类，其中关键的方法是__getitem__（），因为关联的DataLoader依靠这但是Pytorch将__getitem__（）的实现下放到了其他模块中，原因在于不同类型的数据集差异很大，比如图像类数据集

```
class Dataset(Generic[T_co]):
    def __getitem__(self, index) -> T_co:
        raise NotImplementedError

    def __add__(self, other: 'Dataset[T_co]') -> 'ConcatDataset[T_co]':
        return ConcatDataset([self, other])
```

在torchvision中, 可以比较清楚的看到, CIFAR10的数据集继承了VisionDataset (VisionDataset继承了torch.util

```
# torchvision/datasets/cifar.py
```

```
class CIFAR10(VisionDataset):
```

```
    def __init__(
        self,
        root: str,
        train: bool = True,
        transform: Optional[Callable] = None,
        target_transform: Optional[Callable] = None,
        download: bool = False,
    ) -> None:

        super(CIFAR10, self).__init__(root, transform=transform,
                                      target_transform=target_transform)

        #...
        self.data: Any = []
        self.targets = []

        # now load the picked numpy arrays
        for file_name, checksum in downloaded_list:
            file_path = os.path.join(self.root, self.base_folder, file_name)
            with open(file_path, 'rb') as f:
                entry = pickle.load(f, encoding='latin1')
                self.data.append(entry['data'])
                if 'labels' in entry:
                    self.targets.extend(entry['labels'])
                else:
                    self.targets.extend(entry['fine_labels'])

        self.data = np.vstack(self.data).reshape(-1, 3, 32, 32)
        self.data = self.data.transpose((0, 2, 3, 1)) # convert to HWC

        self._load_meta()

    def _load_meta(self) -> None:
        path = os.path.join(self.root, self.base_folder, self.meta['filename'])
        if not check_integrity(path, self.meta['md5']):
            raise RuntimeError('Dataset metadata file not found or corrupted.' +
                               ' You can use download=True to download it')
        with open(path, 'rb') as infile:
            data = pickle.load(infile, encoding='latin1')
            self.classes = data[self.meta['key']]
            self.class_to_idx = {_class: i for i, _class in enumerate(self.classes)}
```

```

def __getitem__(self, index: int) -> Tuple[Any, Any]:
    img, target = self.data[index], self.targets[index]

    img = Image.fromarray(img)

    if self.transform is not None:
        img = self.transform(img)

    if self.target_transform is not None:
        target = self.target_transform(target)

    return img, target

```

从上面的实现中可以看到，CIFAR10的数据集在初始化的时候就把所有的图片都读取并做了初始的处理，考虑到某些我们会有个疑问，所有的worker使用同一个Dataset吗？getitem()会成为瓶颈么？

数据采样

我们在训练模型的时候，一般是把DataLoader当作迭代器来使用，缺省情况下DataLoader只使用一个进程来读取数据，称为_SingleProcessDataLoaderIter，但是当计算速度比较快，比如使用GPU或者多卡进行训练时，为了加快数据加载，可以设置DataLoader使用多进程进行读取，此时DataLoader返回的迭代器称为_MultiProcessingDataLoaderIter。

#Harry torch/utils/data/dataloader.py

```

class DataLoader(Generic[T_co]):
    dataset: Dataset[T_co]
    batch_size: Optional[int]
    num_workers: int
    pin_memory: bool
    drop_last: bool
    timeout: float
    sampler: Union[Sampler, Iterable]
    pin_memory_device: str
    prefetch_factor: int
    _iterator : Optional['_BaseDataLoaderIter']
    __initialized = False

    def _get_iterator(self) -> '_BaseDataLoaderIter':
        if self.num_workers == 0:
            return _SingleProcessDataLoaderIter(self)
        else:
            self.check_worker_number_rationality()
            return _MultiProcessingDataLoaderIter(self)

```

由于要协调进程间数据的读取，_MultiProcessingDataLoaderIter的实现略微复杂一些。首先，在初始化的时候，就multiprocessing库创建多个子进程，每个子进程都在执行_worker_loop()函数。

```
# torch/utils/data/dataloader.py
```

```
class _MultiProcessingDataLoaderIter(_BaseDataLoaderIter):
    def __init__(self, loader):
        #...
        for i in range(self._num_workers):
            index_queue = multiprocessing_context.Queue()
            index_queue.cancel_join_thread()
            w = multiprocessing_context.Process(
                target=_utils.worker._worker_loop,
                args=(self._dataset_kind, self._dataset, index_queue,
                    self._worker_result_queue, self._workers_done_event,
                    self._auto_collation, self._collate_fn, self._drop_last,
                    self._base_seed, self._worker_init_fn, i, self._num_workers,
                    self._persistent_workers, self._shared_seed))
            w.daemon = True
            w.start()
            self._index_queues.append(index_queue)
            self._workers.append(w)
        #...
```

在多进程中环境中，不能使用Python标准库中的Queue。需要使用进程安全的multiprocessing.Queue，和其他语言的进程安全的Queue是_MultiProcessDataLoaderIter中主进程及各个worker子进程之间传递消息的通道，包括以下几种：
- index_queue。存放数据为(send_idx, index)，由main_thread生产，worker_1~n_process消费。其中send_idx是
- worker_result_queue。存放数据为(send_idx, pageble tensor)，由worker_1~n_process产生，pin_memory_thread
- data_queue。存放数据为(send_idx, pinned tensor)，由-pin_memory_thread产生，main_thread消费。

对应于

```
# torch/utils/data/_utils/worker.py
```

```
def _worker_loop(dataset_kind, dataset, index_queue, data_queue, done_event,
                auto_collation, collate_fn, drop_last, base_seed, init_fn, worker_id,
                num_workers, persistent_workers, shared_seed):

    #...
    global _worker_info
    _worker_info = WorkerInfo(id=worker_id, num_workers=num_workers,
                              seed=seed, dataset=dataset)

    #...
    fetcher = _DatasetKind.create_fetcher(dataset_kind, dataset, auto_collation, collate_fn)
    #...
```

```

while watchdog.is_alive():
    r = index_queue.get(timeout=MP_STATUS_CHECK_INTERVAL)

    if isinstance(r, _ResumeIteration):
        #...
        fetcher = _DatasetKind.create_fetcher(
            dataset_kind, dataset, auto_collation, collate_fn, drop_last)
        continue
    #...
    idx, index = r
    data: Union[_IterableDatasetStopIteration, ExceptionWrapper]

    #...
    data = fetcher.fetch(index)
    data_queue.put((idx, data))
    del data, idx, index, r # save memory
    #...

```

这里简单介绍一下fetcher，fetcher的工作就是从Dataset中读取数据，根据Dataset的类型（Map类型或者Iterable）从上面代码可以看出worker的工作流程也比较简单，先根据Dataset类型创建相应的fetcher，然后启动循环，从index中获取数据，值得注意的是，当读到末尾的时候，worker会根据drop_last参数决定是否要丢弃最后这一部分数据，同时如果设置

数据预处理及数据增强

在Dataset的定义中，本身是没有transform参数的，但是我们平时在使用具体的Dataset时，一般都有transform这个参数。

torchvision/datasets/folder.py

```

class DatasetFolder(VisionDataset):
    def __getitem__(self, index: int) -> Tuple[Any, Any]:
        path, target = self.samples[index]
        sample = self.loader(path)
        if self.transform is not None:
            sample = self.transform(sample)
        if self.target_transform is not None:
            target = self.target_transform(target)

        return sample, target

```

锁页内存

到了这里，原始的文件中的数据已经被读取，经过变换后，放到了DataLoader的data_queue里，

```

# torch/utils/data/dataloader.py

class _MultiProcessingDataLoaderIter(_BaseDataLoaderIter):
    def __init__(self, loader):
        #...

        if self._pin_memory:
            self._pin_memory_thread_done_event = threading.Event()

            # Queue is not type-annotated
            self._data_queue = queue.Queue() # type: ignore[var-annotated]
            pin_memory_thread = threading.Thread(
                target=_utils.pin_memory._pin_memory_loop,
                args=(self._worker_result_queue, self._data_queue,
                    torch.cuda.current_device(),
                    self._pin_memory_thread_done_event, self._pin_memory_device))
            pin_memory_thread.daemon = True
            pin_memory_thread.start()
            # Similar to workers (see comment above), we only register
            # pin_memory_thread once it is started.
            self._pin_memory_thread = pin_memory_thread
        else:
            self._data_queue = self._worker_result_queue

```

可以看到，DataLoader只启动了一个pin_memory的线程，这个线程的工作相当简单，就是将_data_queue中的样本数
 TODO: Tensor的pin_memory方法以后有机会可以再看一下。

数据加载到GPU

数据分发

DistributedSampler

torch/utils/dataset.py

模型训练中的数据集

—

设计原则1. DataLoader -> Dataset

参考

- 万字综述，核心开发者全面解读PyTorch内部机制 <https://zhuanlan.zhihu.com/p/67834038>
- <https://blog.csdn.net/u013608424/article/details/123782284> # 第9章
优化器

分布式

本章主要内容

- 为什么需要分布式
- 分布式的难点在哪里？
- PyTorch中的相关模块
 - THD
 - C10D
 - torch.multiprocessing
 - torch.distributedDataParallel (DP)
 - DistributedDataParallel (DDP)
 - torch.distributed.rpc

什么是分布式训练

分布式计算

由于单个节点的计算能力有限，对于计算密集型的任务，只在单个节点上运行，可能会花费非常多的时间，此时充分将任务从单节点转化为分布式任务，需要考虑不同节点间的通信，包括输入数据的拆分，临时数据的分发与归并，计

为了简化算法开发的复杂度，将分布式计算中的数据分发和网络通信与具体的算法应用分开，先驱们开发了不同的分

在深度学习领域，模型的效果主要来自于两个方面：海量的数据和精心设计的复杂网络结构，这两点使得深度学习机

来源：Compute Trends Across Three Eras of Machine Learning

深度学习模型分布式训练的进展

PyTorch中的分布式训练

参考

- <https://zhuanlan.zhihu.com/p/136372142>

第11章 JIT

TorchScript

为什么需要JIT

- 性能

实现JIT的挑战

- 动态图中的条件逻辑

一个简单的例子

为了说明JIT是如何工作的，我们看一个简单的例子：

```
@torch.jit.script
def foo(len):
    # type: (int) -> torch.Tensor
    rv = torch.zeros(3, 4)
    for i in range(len):
        if i < 10:
            rv = rv - 1.0
        else:
            rv = rv + 1.0
    return rv
```

```
print(foo.code)
```

加上修饰器后，上面的函数foo的类型变成了，并且其代码被重新编译成了下面的形式：

```
def foo(len: int) -> Tensor:
    rv = torch.zeros([3, 4], dtype=None, layout=None, device=None, pin_memory=None)
    rv0 = rv
    for i in range(len):
        if torch.lt(i, 10):
            rv1 = torch.sub(rv0, 1., 1)
        else:
            rv1 = torch.add(rv0, 1., 1)
        rv0 = rv1
    return rv0
```

可见其中基本的条件语句被转换成了torch的函数，但这仍然是Python代码层面，在执行层，TorchScript使用的是静态单赋值（SSA）中间表示（IR），其中的指令包括ATen（the C++ backend of PyTorch）算子及其他一些原语，比如条件控制和循环控制的原语。

如果我们打印`print(foo.graph)`，可以看到如下的输出，其中“:5:4”这样的注释代表中间代码所对应的Python源Notebook，读者朋友可以忽略文件名，只关注代码位置即可。

```
graph(%len.1 : int):
  %20 : int = prim::Constant[value=1]()
  %13 : bool = prim::Constant[value=1]() # <ipython-input-4-01a58e79a588>:5:4
  %5 : None = prim::Constant()
  %1 : int = prim::Constant[value=3]() # <ipython-input-4-01a58e79a588>:4:21
  %2 : int = prim::Constant[value=4]() # <ipython-input-4-01a58e79a588>:4:24
  %16 : int = prim::Constant[value=10]() # <ipython-input-4-01a58e79a588>:6:15
  %19 : float = prim::Constant[value=1]() # <ipython-input-4-01a58e79a588>:7:22
  %4 : int[] = prim::ListConstruct(%1, %2)
  %rv.1 : Tensor = aten::zeros(%4, %5, %5, %5, %5) # <ipython-input-4-01a58e79a588>:4:9
  %rv : Tensor = prim::Loop(%len.1, %13, %rv.1) # <ipython-input-4-01a58e79a588>:5:4
  block0(%i.1 : int, %rv.14 : Tensor):
    %17 : bool = aten::lt(%i.1, %16) # <ipython-input-4-01a58e79a588>:6:11
    %rv.13 : Tensor = prim::If(%17) # <ipython-input-4-01a58e79a588>:6:8
    block0():
      %rv.3 : Tensor = aten::sub(%rv.14, %19, %20) # <ipython-input-4-01a58e79a588>:7:1
      -> (%rv.3)
    block1():
      %rv.6 : Tensor = aten::add(%rv.14, %19, %20) # <ipython-input-4-01a58e79a588>:9:1
      -> (%rv.6)
    -> (%13, %rv.13)
  return (%rv)
```

JIT trace的实现

```
def fill_row_zero(x):
    x[0] = torch.rand(*x.shape[1:2])
    return x

traced = torch.jit.trace(fill_row_zero, (torch.rand(3, 4),))
print(traced.graph)
```

Trace的实现在这里（不同版本的实现位置可能不一样）：

```
# torch/jit/_trace.py

def trace(
    func,
    example_inputs,
    optimize=None,
    check_trace=True,
    check_inputs=None,
    check_tolerance=1e-5,
```

```

    strict=True,
    _force_outplace=False,
    _module_class=None,
    _compilation_unit=_python_cu,
):

    #YL      Module   trace_module

    var_lookup_fn = _create_interpreter_name_lookup_fn(0)

    name = _qualified_name(func)
    traced = torch._C._create_function_from_trace(
        name,
        func,
        example_inputs,
        var_lookup_fn,
        strict,
        _force_outplace,
        get_callable_argument_names(func)
    )

    # Check the trace against new traces created from user-specified inputs

    return traced

```

_C是torch的C++模块，因此该调用转到了C++部分，在初始化的时候，_create_function_from_trace被注册到了torch的C++模块中。
//YL torch/csrc/jit/python/script_init.cpp

```

m.def(
    "_create_function_from_trace",
    [](const std::string& qualname,
        const py::function& func,
        const py::tuple& input_tuple,
        const py::function& var_name_lookup_fn,
        bool strict,
        bool force_outplace,
        const std::vector<std::string>& argument_names) {
        auto typed_inputs = toTraceableStack(input_tuple);
        std::shared_ptr<Graph> graph = std::get<0>(tracer::createGraphByTracing(
            func,
            typed_inputs,
            var_name_lookup_fn,
            strict,
            force_outplace,
            /*self=*/nullptr,
            argument_names));
    }
);

```

```

    auto cu = get_python_cu();
    auto name = c10::QualifiedName(qualname);
    auto result = cu->create_function(
        std::move(name), std::move(graph), /*shouldMangle=*/true);
    StrongFunctionPtr ret(std::move(cu), result);
    didFinishEmitFunction(ret);
    return ret;
},
py::arg("name"),
py::arg("func"),
py::arg("input_tuple"),
py::arg("var_name_lookup_fn"),
py::arg("strict"),
py::arg("force_outplace"),
py::arg("argument_names") = std::vector<std::string>());

```

可以看到，主要的工作是构造一个Graph，并且是由`tracer::createGraphByTracing()`完成的。

参考

- <https://pytorch.org/docs/stable/jit.html>
- <https://zhuanlan.zhihu.com/p/410507557>

第3章 自动微分

Index

- 理论知识
- 梯度的保存
- 梯度的计算
- 反向传播

梯度的初步认识

我们知道，深度神经网络的训练时依赖于梯度的反向传播的，因此在深度学习框架的设计上就涉及到几个问题：

- 梯度保存在哪里？
- 梯度是怎样计算的？
- 神经网络的参数是如何更新的？
- 如何实现反向传播？

神经网络的核心数据结构是Tensor，对于需要优化的Tensor，每次更新，都会有一个对应的梯度。因此最合适的方式是在初始化Tensor的时候，可以指定一个参数`requires_grad`，代表这个Tensor是否需要计算梯度。在涉及复杂的神经网络之前，我们先看一个非常简单的计算，这个例子来自于pytorch官方文档。

```
import torch
```

```
x = torch.ones(2, 2, requires_grad=True)
print(x)
```

输出结果为:

```
tensor([[1., 1.],
        [1., 1.]], requires_grad=True)
```

如果对这个Tensor做一些操作:

```
y = x + 2
print(y)
```

输出为:

```
tensor([[3., 3.],
        [3., 3.]], grad_fn=<AddBackward0>)
```

可以看到基于加法操作的Tensor y, 被附加了一个grad_fn的函数。因为x是需要梯度的, 而y是基于x的加法操作得到的。同理做更多的操作:

```
z = y * y * 3
out = z.mean()
```

```
print(z, out)
```

输出如下, 可见计算梯度的函数不是固定的, 不同的操作对应不同的梯度计算函数。

```
tensor([[27., 27.],
        [27., 27.]], grad_fn=<MulBackward0>)
tensor(27., grad_fn=<MeanBackward0>)
```

现在我们再看一下梯度的计算和反向传播过程, 刚才提到梯度是保存在Tensor里的, 在pytorch中, 可以通过Tensor的backward()方法。

```
print(x.grad)
```

输出:

```
tensor([[4.5000, 4.5000],
        [4.5000, 4.5000]])
```

关于梯度的基本理论

雅克比矩阵

一元Tensor的梯度计算，不需要雅克比矩阵

待补充

PyTorch中梯度的计算过程

从刚才的例子可以看到，梯度可以通过Tensor.backward()函数计算得到。那么这个函数都做了什么呢？

```
class Tensor(torch._C._TensorBase):
    def backward(self, gradient=None, retain_graph=None, create_graph=False, inputs=None):

        if has_torch_function_unary(self):
            return handle_torch_function(
                Tensor.backward,
                (self,),
                self,
                gradient=gradient,
                retain_graph=retain_graph,
                create_graph=create_graph,
                inputs=inputs)

        torch.autograd.backward(self, gradient, retain_graph, create_graph, inputs=inputs)
```

我们先忽略对一元情况的处理，一般来说，最终会调用autograd.backward()函数进行梯度的计算，这个函数定义在这个函数在计算梯度并且反向传播的时候，会把梯度保存在计算图的叶子节点中。需要注意的是，在调用backward前

```
def backward(
    tensors: _TensorOrTensors,
    grad_tensors: Optional[_TensorOrTensors] = None,
    retain_graph: Optional[bool] = None,
    create_graph: bool = False,
    grad_variables: Optional[_TensorOrTensors] = None,
    inputs: Optional[_TensorOrTensors] = None,
) -> None:
    if grad_variables is not None:
        warnings.warn("'grad_variables' is deprecated. Use 'grad_tensors' instead.")
    if grad_tensors is None:
        grad_tensors = grad_variables
    else:
        raise RuntimeError("'grad_tensors' and 'grad_variables' (deprecated) "
                           "arguments both passed to backward(). Please only "
                           "use 'grad_tensors'.")
    if inputs is not None and len(inputs) == 0:
```

```

        raise RuntimeError("'inputs' argument to backward() cannot be empty.")

    tensors = (tensors,) if isinstance(tensors, torch.Tensor) else tuple(tensors)
    inputs = (inputs,) if isinstance(inputs, torch.Tensor) else \
        tuple(inputs) if inputs is not None else tuple()

    grad_tensors_ = _tensor_or_tensors_to_tuple(grad_tensors, len(tensors))
    grad_tensors_ = _make_grads(tensors, grad_tensors_, is_grads_batched=False)
    if retain_graph is None:
        retain_graph = create_graph

    # The reason we repeat same the comment below is that
    # some Python versions print out the first line of a multi-line function
    # calls in the traceback and some print out the last line
    Variable._execution_engine.run_backward( # Calls into the C++ engine to run the backward
        tensors, grad_tensors_, retain_graph, create_graph, inputs,
        allow_unreachable=True, accumulate_grad=True) # Calls into the C++ engine to run the

```

在经过一些处理之后，最后调用的是Variable._execution_engine.run_backward()函数，但事实上，Variable._exe

```

import torch
from torch._six import with_metaclass

class VariableMeta(type):
    def __instancecheck__(cls, other):
        return isinstance(other, torch.Tensor)

# mypy doesn't understand torch._six.with_metaclass
class Variable(with_metaclass(VariableMeta, torch._C._LegacyVariableBase)): # type: ignore
    pass

from torch._C import _ImperativeEngine as ImperativeEngine
Variable._execution_engine = ImperativeEngine()

```

在对应的C++代码中，使用PyModule_AddObject注册了_ImperativeEngine这个类对象。
torch/csrc/autograd/python_engine.cpp

```

PyTypeObject THPEngineType = {
    PyVarObject_HEAD_INIT(nullptr, 0) "torch._C._EngineBase", /* tp_name */
    sizeof(THPEngine), /* tp_basicsize */
    0, /* tp_itemsize */
    nullptr, /* tp_dealloc */
    0, /* tp_vectorcall_offset */
    nullptr, /* tp_getattr */
    nullptr, /* tp_setattr */
    nullptr, /* tp_reserved */
    nullptr, /* tp_repr */

```



```

    nullptr, /* tp_as_number */
    nullptr, /* tp_as_sequence */
    nullptr, /* tp_as_mapping */
    nullptr, /* tp_hash */
    nullptr, /* tp_call */
    nullptr, /* tp_str */
    nullptr, /* tp_getattro */
    nullptr, /* tp_setattro */
    nullptr, /* tp_as_buffer */
    Py_TPFLAGS_DEFAULT | Py_TPFLAGS_BASETYPE, /* tp_flags */
    nullptr, /* tp_doc */
    nullptr, /* tp_traverse */
    nullptr, /* tp_clear */
    nullptr, /* tp_richcompare */
    0, /* tp_weaklistoffset */
    nullptr, /* tp_iter */
    nullptr, /* tp_iternext */
    THPEngine_methods, /* tp_methods */
    nullptr, /* tp_members */
    nullptr, /* tp_getset */
    nullptr, /* tp_base */
    nullptr, /* tp_dict */
    nullptr, /* tp_descr_get */
    nullptr, /* tp_descr_set */
    0, /* tp_dictoffset */
    nullptr, /* tp_init */
    nullptr, /* tp_alloc */
    THPEngine_new /* tp_new */
};

bool THPEngine_initModule(PyObject* module) {
#ifdef _WIN32
    if (pthread_atfork(nullptr, nullptr, child_atfork) != 0) {
        throw std::runtime_error("unable to set pthread_atfork handler");
    }
#endif
    if (PyType_Ready(&THPEngineType) < 0)
        return false;
    Py_INCREF(&THPEngineType);
    PyModule_AddObject(module, "_ImperativeEngine", (PyObject*)&THPEngineType);
    set_default_engine_stub(python::PythonEngine::get_python_engine());
    return true;
}

```

希望了解PyModule_AddObject细节的同学可以学习一下Cython。在这里我们只需要知道这个函数可以将C++的类型注册到Python中。可以看到，实际注册的对象是一个PyTypeObject。PyTypeObject是Python中非常重要的一种类型，PyTypeObject就是

参考 <https://blog.csdn.net/zhangyifei216/article/details/50581787>

对象中每个字段的含义可以从注释中看出来，不过基本可以忽略，大部分都是空，最后一个字段是THPEngine_new，有一点待确认，就是PyTypeObject各个字段的定义，在不同Python版本中估计是不一样的，如何保证兼容呢？至少对于_ImperativeEngine这个类，在C++中注册了以下几个函数，其中就包括run_backward函数，对应的C++实现是TH

```
// NOLINTNEXTLINE(cppcoreguidelines-avoid-c-arrays,modernize-avoid-c-arrays,cppcoreguidelines-avoid-const-refs-in-large-structs)
static struct PyMethodDef THPEngine_methods[] = {
    {(char*)"run_backward",
     castPyCFunctionWithKeywords(THPEngine_run_backward),
     METH_VARARGS | METH_KEYWORDS,
     nullptr},
    {(char*)"queue_callback", THPEngine_queue_callback, METH_0, nullptr},
    {(char*)"is_checkpoint_valid",
     THPEngine_is_checkpoint_valid,
     METH_NOARGS,
     nullptr},
    {nullptr}};
```

THPEngine_run_backward函数的实现相对比较复杂，但是其中开始部分是对输入参数进行解析，在结束部分是对Ten

```
// Implementation of torch._C._EngineBase.run_backward
PyObject* THPEngine_run_backward(
    PyObject* self,
    PyObject* args,
    PyObject* kwargs) {

    HANDLE_TH_ERRORS
    PyObject* tensors = nullptr;
    PyObject* grad_tensors = nullptr;
    unsigned char keep_graph = 0;
    unsigned char create_graph = 0;
    PyObject* inputs = nullptr;
    unsigned char allow_unreachable = 0;
    unsigned char accumulate_grad =
        0; // Indicate whether to accumulate grad into leaf Tensors or capture
    const char* accepted_kwargs[] = {// NOLINT
        "tensors",
        "grad_tensors",
        "keep_graph",
        "create_graph",
        "inputs",
        "allow_unreachable",
        "accumulate_grad",
        nullptr};

    if (!PyArg_ParseTupleAndKeywords(
```

```

        args,
        kwargs,
        "00bb|0bb",
        (char**)accepted_kwargs,
        &tensors,
        &grad_tensors,
        &keep_graph,
        &create_graph,
        &inputs,
        &allow_unreachable,
        &accumulate_grad))
    return nullptr;

// ... check arguments

// ... init edges

variable_list outputs;
{
    pybind11::gil_scoped_release no_gil;
    auto& engine = python::PythonEngine::get_python_engine();
    outputs = engine.execute(
        roots, grads, keep_graph, create_graph, accumulate_grad, output_edges);
}

// ... assign gradients to Tensor
}

```

在执行run_backward()函数时，首先通过PyArg_ParseTupleAndKeywords()函数对入参进行格式解析，将Python的对
可以看到，计算梯度的核心函数是engine.execute()，PythonEngine继承自Engine，实现execute()的时候也是简单
下面的代码来自于torch/csrc/autograd/python_engine.h 和torch/csrc/autograd/python_engine.cpp。

```

struct PythonEngine : public Engine {
    static Engine& get_python_engine();
    ~PythonEngine() override;
    void thread_init(
        int device,
        const std::shared_ptr<ReadyQueue>& ready_queue,
        bool should_increment) override;
    void thread_on_exception(
        std::shared_ptr<GraphTask> graph_task,
        const std::shared_ptr<Node>& fn,
        std::exception& e) override;
    variable_list execute(
        const edge_list& roots,

```

```

        const variable_list& inputs,
        bool keep_graph,
        bool create_graph,
        bool accumulate_grad,
        const edge_list& outputs = {}) override;

c10::intrusive_ptr<at::ivalue::Future> execute_with_graph_task(
    const std::shared_ptr<GraphTask>& graph_task,
    std::shared_ptr<Node> graph_root,
    InputBuffer&& input_buffer) override;

std::unique_ptr<AnomalyMetadata> make_anomaly_metadata() override;
std::unique_ptr<SavedVariableHooks> get_default_saved_variable_hooks()
    override;

private:
    PythonEngine();
};

Engine& PythonEngine::get_python_engine() {
    static PythonEngine engine;
    // This is "probably" thread-safe because the flag is set in a fork handler
    // before any threads are created, and this function is only called with the
    // GIL held. However, using fork + threads is playing with fire so this is
    // more of a "best effort" thing. For example, if the fork occurs while the
    // backwards threads hold a lock, we'll probably deadlock in the engine
    // destructor.
    if (_reinitialize_engine) {
        engine.release_workers();
        engine.~PythonEngine();
        new (&engine) torch::autograd::python::PythonEngine();
        _reinitialize_engine = false;
    }
    return engine;
}

variable_list PythonEngine::execute(
    const edge_list& roots,
    const variable_list& inputs,
    bool keep_graph,
    bool create_graph,
    bool accumulate_grad,
    const edge_list& outputs) {
    TORCH_CHECK(
        !PyGILState_Check(),
        "The autograd engine was called while holding the GIL. If you are using the C++ "

```

```

        "API, the autograd engine is an expensive operation that does not require the "
        "GIL to be held so you should release it with 'pybind11::gil_scoped_release no_gil;'"
        ". If you are not using the C++ API, please report a bug to the pytorch team.")
    try {
        return Engine::execute(
            roots, inputs, keep_graph, create_graph, accumulate_grad, outputs);
    } catch (python_error& e) {
        e.restore();
        throw;
    }
}

```

Engine的定义和实现分别在torch/csrc/autograd/engine.h和torch/csrc/autograd/engine.cpp中。

在一个平台级的系统里，能够被命名为Engine的类型，一定是整个系统的核心，而Engine.execute()函数的实现肯定是这个核心对象的主要执行逻辑，在深度学习框架中，这个最主要的执行逻辑就是

```

auto Engine::execute(
    const edge_list& roots,
    const variable_list& inputs,
    bool keep_graph,
    bool create_graph,
    bool accumulate_grad,
    const edge_list& outputs) -> variable_list {
    // NOLINTNEXTLINE(cppcoreguidelines-pro-type-const-cast)
    validate_outputs(
        roots, const_cast<variable_list&>(inputs), [](const std::string& msg) {
            return msg;
        });
    if (accumulate_grad && create_graph) {
        TORCH_WARN_ONCE(
            "Using backward() with create_graph=True will create a reference cycle "
            "between the parameter and its gradient which can cause a memory leak. "
            "We recommend using autograd.grad when creating the graph to avoid this. "
            "If you have to use this function, make sure to reset the .grad fields of "
            "your parameters to None after use to break the cycle and avoid the leak.");
    }

    // accumulate_grad is true if and only if the frontend call was to
    // grad(), not backward(). grad() returns the sum of the gradients
    // w.r.t. the inputs and thus needs the inputs to be present.
    TORCH_CHECK_VALUE(
        accumulate_grad || !outputs.empty(), "grad requires non-empty inputs.");

    // A fresh first time Engine::execute call should start on the CPU device,
    // initialize a new thread local ready queue on CPU or reuse the existing one
    // (if there is one allocated already, i.e. consecutive backward calls,

```

```

// re-entrant backward calls), then memoize the local_ready_queue in GraphTask
init_local_ready_queue();
bool not_reentrant_backward_call = worker_device == NO_DEVICE;

auto graph_task = std::make_shared<GraphTask>(
    /* keep_graph */ keep_graph,
    /* create_graph */ create_graph,
    /* depth */ not_reentrant_backward_call ? 0 : total_depth + 1,
    /* cpu_ready_queue */ local_ready_queue);

// If we receive a single root, skip creating extra root node
bool skip_dummy_node = roots.size() == 1;
auto graph_root = skip_dummy_node
    ? roots.at(0).function
    : std::make_shared<GraphRoot>(roots, inputs);

auto min_topo_nr = compute_min_topological_nr(outputs);
// Now compute the dependencies for all executable functions
compute_dependencies(graph_root.get(), *graph_task, min_topo_nr);

if (!outputs.empty()) {
    graph_task->init_to_execute(
        *graph_root, outputs, accumulate_grad, min_topo_nr);
}

// Queue the root
if (skip_dummy_node) {
    InputBuffer input_buffer(roots.at(0).function->num_inputs());
    auto input = inputs.at(0);

    const auto input_stream = InputMetadata(input).stream();
    const auto opt_next_stream =
        roots.at(0).function->stream(c10::DeviceType::CUDA);
    input_buffer.add(
        roots.at(0).input_nr, std::move(input), input_stream, opt_next_stream);

    execute_with_graph_task(graph_task, graph_root, std::move(input_buffer));
} else {
    execute_with_graph_task(
        graph_task, graph_root, InputBuffer(variable_list()));
}

// Avoid a refcount bump for the Future, since we check for refcount in
// DistEngine (see TORCH_INTERNAL_ASSERT(futureGrads.use_count() == 1)
// in dist_engine.cpp).
auto& fut = graph_task->future_result_;
fut->wait();

```

```

graph_task->warning_handler_.replay_warnings();
return fut->value().toTensorVector();
}

```

GraphTask在执行的过程中创建出来的。

明显能够看出，execute()方法中的重要步骤是execute_with_graph_task()函数。

执行的时候就是对graph_task进行BFS遍历，从root开始调用各Node的operator()重载函数。

```

c10::intrusive_ptr<at::IValue::Future> Engine::execute_with_graph_task(
    const std::shared_ptr<GraphTask>& graph_task,
    std::shared_ptr<Node> graph_root,
    InputBuffer&& input_buffer) {
    initialize_device_threads_pool();
    // Lock mutex for GraphTask.
    std::unique_lock<std::mutex> lock(graph_task->mutex_);

    auto queue = ready_queue(graph_task->cpu_ready_queue_, input_buffer.device());

    // worker_device == NO_DEVICE it's a CPU thread and it's trying to drive the
    // autograd engine with corresponding GraphTask, and its NOT a re-entrant call
    if (worker_device == NO_DEVICE) {
        // We set the worker_device to CPU_DEVICE only if worker_device was
        // previously NO_DEVICE. Setting it to CPU afterwards allow us to detect
        // whether this is a re-entrant call or not.
        set_device(CPU_DEVICE);

        // set the graph_task owner to the current device
        graph_task->owner_ = worker_device;

        // Now that all the non-thread safe fields of the graph_task have been
        // populated, we can enqueue it.
        queue->push(
            NodeTask(graph_task, std::move(graph_root), std::move(input_buffer)));

        // The owning thread start to drive the engine execution for any CPU task
        // that was just pushed or will be added later from other worker threads
        lock.unlock();
        thread_main(graph_task);
        TORCH_INTERNAL_ASSERT(graph_task->future_result_->completed());
        // reset the worker_device after the completion of the graph_task, this is
        // so that the initial state of the engine remains the same across every
        // backward() or grad() call, we don't need to reset local_ready_queue as we
        // could possibly reuse it for new backward calls.
        worker_device = NO_DEVICE;
    } else {
        // If worker_device is any devices (i.e. CPU, CUDA): this is a re-entrant

```

```

//    backward call from that device.
graph_task->owner_ = worker_device;

// Now that all the non-thread safe fields of the graph_task have been
// populated, we can enqueue it.
queue->push(
    NodeTask(graph_task, std::move(graph_root), std::move(input_buffer)));

if (current_depth >= max_recursion_depth_) {
    // See Note [Reentrant backwards]
    // If reached the max depth, switch to a different thread
    add_thread_pool_task(graph_task);
} else {
    // Total depth needs to be updated only in this codepath, since it is
    // not used in the block above (when we call add_thread_pool_task).
    // In the codepath above, GraphTask.reentrant_depth_ is used to
    // bootstrap total_depth in the other thread.
    ++total_depth;

    // Get back to work while we wait for our new graph_task to
    // complete!
    ++current_depth;
    lock.unlock();
    thread_main(graph_task);
    --current_depth;
    --total_depth;

    // The graph task should have completed and the associated future should
    // be marked completed as well since 'thread_main' above is a call
    // blocking an autograd engine thread.
    TORCH_INTERNAL_ASSERT(graph_task->future_result_->completed());
}
}
// graph_task_exec_post_processing is done when the Future is marked as
// completed in mark_as_completed_and_run_post_processing.
return graph_task->future_result_;
}

```

这里涉及到几个逻辑： - 梯度的计算一般也是矩阵计算，对算力要求比较高，在有GPU的情况下可以使用GPU计算，且由于计算图是一个有向无环图，计算的时候有很多可以并行的节点，因此在设计上可以将任务推到队列中进行并行

从上面的代码可以看到，计算的核心是thread_main(graph_task)

```

auto Engine::thread_main(const std::shared_ptr<GraphTask>& graph_task) -> void {
    // When graph_task is nullptr, this is a long running thread that processes
    // tasks (ex: device threads). When graph_task is non-null (ex: reentrant
    // backwards, user thread), this function is expected to exit once that

```



```

    // graph_task complete.

#ifdef USE_ROCM
    // Keep track of backward pass for rocblas.
    at::ROCMBackwardPassGuard in_backward;
#endif

    // local_ready_queue should already been initialized when we get into
    // thread_main
    TORCH_INTERNAL_ASSERT(local_ready_queue != nullptr);
    while (graph_task == nullptr || !graph_task->future_result_->completed()) {
        // local_graph_task represents the graph_task we retrieve from the queue.
        // The outer graph_task represents the overall graph_task we need to execute
        // for reentrant execution.
        std::shared_ptr<GraphTask> local_graph_task;
        {
            // Scope this block of execution since NodeTask is not needed after this
            // block and can be deallocated (release any references to grad tensors
            // as part of inputs_).
            NodeTask task = local_ready_queue->pop();
            // This will only work if the worker is running a non backward task
            // TODO Needs to be fixed this to work in all cases
            if (task.isShutdownTask_) {
                C10_LOG_API_USAGE_ONCE("torch.autograd.thread_shutdown");
                break;
            }

            if (!(local_graph_task = task.base_.lock())) {
                // GraphTask for function is no longer valid, skipping further
                // execution.
                continue;
            }

            if (task.fn_ && !local_graph_task->has_error_.load()) {
                // Set the ThreadLocalState before calling the function.
                // NB: The ThreadLocalStateGuard doesn't set the grad_mode because
                // GraphTask always saves ThreadLocalState without grad_mode.
                at::ThreadLocalStateGuard tls_guard(local_graph_task->thread_locals_);
                c10::Warning::WarningHandlerGuard warnings_guard(
                    &local_graph_task->warning_handler_);

                try {
                    // The guard sets the thread_local current_graph_task on construction
                    // and restores it on exit. The current_graph_task variable helps
                    // queue_callback() to find the target GraphTask to append final
                    // callbacks.

```

```

GraphTaskGuard guard(local_graph_task);
NodeGuard ndguard(task.fn_);
{
    RECORD_FUNCTION(
        c10::str(
            "autograd::engine::evaluate_function: ",
            task.fn_.get()->name()),
        c10::ArrayRef<const c10::IValue>());
    evaluate_function(
        local_graph_task,
        task.fn_.get(),
        task.inputs_,
        local_graph_task->cpu_ready_queue_);
}
} catch (std::exception& e) {
    thread_on_exception(local_graph_task, task.fn_, e);
}
}
}

// Decrement the outstanding tasks.
--local_graph_task->outstanding_tasks_;

// Check if we've completed execution.
if (local_graph_task->completed()) {
    local_graph_task->mark_as_completed_and_run_post_processing();

    auto base_owner = local_graph_task->owner_;
    // The current worker thread finish the graph_task, but the owning thread
    // of the graph_task might be sleeping on pop() if it does not have work.
    // So we need to send a dummy function task to the owning thread just to
    // ensure that it's not sleeping, so that we can exit the thread_main.
    // If it has work, it might see that graph_task->outstanding_tasks_ == 0
    // before it gets to the task, but it's a no-op anyway.
    //
    // NB: This is not necessary if the current thread is the owning thread.
    if (worker_device != base_owner) {
        // Synchronize outstanding_tasks_ with queue mutex
        std::atomic_thread_fence(std::memory_order_release);
        ready_queue_by_index(local_graph_task->cpu_ready_queue_, base_owner)
            ->push(NodeTask(local_graph_task, nullptr, InputBuffer(0)));
    }
}
}
}
}

```

thread_main()方法的最重要的步骤是调用evaluate_function()。

```
void Engine::evaluate_function(
    std::shared_ptr<GraphTask>& graph_task,
    Node* func,
    InputBuffer& inputs,
    const std::shared_ptr<ReadyQueue>& cpu_ready_queue) {
    // The InputBuffer::adds that supplied incoming grads took pains to
    // ensure they're safe to consume in the context of the present
    // func's stream (if applicable). So we guard onto that stream
    // before working with the grads in any capacity.
    const auto opt_parent_stream = (*func).stream(c10::DeviceType::CUDA);
    c10::OptionalStreamGuard parent_stream_guard{opt_parent_stream};

    // If exec_info_ is not empty, we have to instrument the execution
    auto& exec_info_ = graph_task->exec_info_;
    if (!exec_info_.empty()) {
        auto& fn_info = exec_info_.at(func);
        if (auto* capture_vec = fn_info.captures_.get()) {
            // Lock mutex for writing to graph_task->captured_vars_.
            std::lock_guard<std::mutex> lock(graph_task->mutex_);
            for (const auto& capture : *capture_vec) {
                auto& captured_grad = graph_task->captured_vars_[capture.output_idx_];
                captured_grad = inputs[capture.input_idx_];
                for (auto& hook : capture.hooks_) {
                    captured_grad = (*hook)(captured_grad);
                }
                if (opt_parent_stream) {
                    // No need to take graph_task->mutex_ here, we already hold it
                    graph_task->leaf_streams.emplace(*opt_parent_stream);
                }
            }
        }
        if (!fn_info.needed_) {
            // Skip execution if we don't need to execute the function.
            return;
        }
    }

    auto outputs = call_function(graph_task, func, inputs);

    auto& fn = *func;
    if (!graph_task->keep_graph_) {
        fn.release_variables();
    }
}
```

```

int num_outputs = outputs.size();
if (num_outputs == 0) { // Note: doesn't acquire the mutex
    // Records leaf stream (if applicable)
    // See Note [Streaming backwards]
    if (opt_parent_stream) {
        std::lock_guard<std::mutex> lock(graph_task->mutex_);
        graph_task->leaf_streams.emplace(*opt_parent_stream);
    }
    return;
}

if (AnomalyMode::is_enabled()) {
    AutoGradMode grad_mode(false);
    for (const auto i : c10::irange(num_outputs)) {
        auto& output = outputs[i];
        at::OptionalDeviceGuard guard(device_of(output));
        if (output.defined() && isnan(output).any().item<uint8_t>()) {
            std::stringstream ss;
            ss << "Function '" << fn.name() << "' returned nan values in its " << i
                << "th output.";
            throw std::runtime_error(ss.str());
        }
    }
}

// Lock mutex for the accesses to GraphTask dependencies_, not_ready_ and
// cpu_ready_queue_ below
std::lock_guard<std::mutex> lock(graph_task->mutex_);
for (const auto i : c10::irange(num_outputs)) {
    auto& output = outputs[i];
    const auto& next = fn.next_edge(i);

    if (!next.is_valid())
        continue;

    // Check if the next function is ready to be computed
    bool is_ready = false;
    auto& dependencies = graph_task->dependencies_;
    auto it = dependencies.find(next.function.get());

    if (it == dependencies.end()) {
        auto name = next.function->name();
        throw std::runtime_error(string("dependency not found for ") + name);
    } else if (--it->second == 0) {
        dependencies.erase(it);
        is_ready = true;
    }
}

```

```

    }

    auto& not_ready = graph_task->not_ready_;
    auto not_ready_it = not_ready.find(next.function.get());
    if (not_ready_it == not_ready.end()) {
        // Skip functions that aren't supposed to be executed
        if (!exec_info_.empty()) {
            auto it = exec_info_.find(next.function.get());
            if (it == exec_info_.end() || !it->second.should_execute()) {
                continue;
            }
        }
        // No buffers have been allocated for the function
        InputBuffer input_buffer(next.function->num_inputs());

        // Accumulates into buffer
        const auto opt_next_stream = next.function->stream(c10::DeviceType::CUDA);
        input_buffer.add(
            next.input_nr, std::move(output), opt_parent_stream, opt_next_stream);

        if (is_ready) {
            auto queue = ready_queue(cpu_ready_queue, input_buffer.device());
            queue->push(
                NodeTask(graph_task, next.function, std::move(input_buffer)));
        } else {
            not_ready.emplace(next.function.get(), std::move(input_buffer));
        }
    } else {
        // The function already has a buffer
        auto& input_buffer = not_ready_it->second;

        // Accumulates into buffer
        const auto opt_next_stream = next.function->stream(c10::DeviceType::CUDA);
        input_buffer.add(
            next.input_nr, std::move(output), opt_parent_stream, opt_next_stream);
        if (is_ready) {
            auto queue = ready_queue(cpu_ready_queue, input_buffer.device());
            queue->push(
                NodeTask(graph_task, next.function, std::move(input_buffer)));
            not_ready.erase(not_ready_it);
        }
    }
}
}
}

```

其核心操作是这一个调用：

```

auto outputs = call_function(graph_task, func, inputs);
call_function的实现也在engine.cpp中。

static variable_list call_function(
    std::shared_ptr<GraphTask>& graph_task,
    Node* func,
    InputBuffer& inputBuffer) {
    CheckpointValidGuard cpvguard(graph_task);
    auto& fn = *func;
    auto inputs =
        call_pre_hooks(fn, InputBuffer::variables(std::move(inputBuffer)));

    if (!graph_task->keep_graph_) {
        fn.will_release_variables();
    }

    const auto has_post_hooks = !fn.post_hooks().empty();
    variable_list outputs;

    if (has_post_hooks) {
        // In functions/accumulate_grad.cpp, there is some logic to check the
        // conditions under which the incoming gradient can be stolen directly
        // (which elides a deep copy) instead of cloned. One of these conditions
        // is that the incoming gradient's refcount must be 1 (nothing else is
        // referencing the same data). Stashing inputs_copy here bumps the
        // refcount, so if post hooks are employed, it's actually still ok for
        // accumulate_grad.cpp to steal the gradient if the refcount is 2.
        //
        // "new_grad.use_count() <= 1 + !post_hooks().empty()" in
        // accumulate_grad.cpp accounts for this, but also creates a silent
        // dependency between engine.cpp (ie, this particular engine
        // implementation) and accumulate_grad.cpp.
        //
        // If you change the logic here, make sure it's compatible with
        // accumulate_grad.cpp.
        auto inputs_copy = inputs;
        outputs = fn(std::move(inputs_copy));
    } else {
        outputs = fn(std::move(inputs));
    }

    validate_outputs(fn.next_edges(), outputs, [&](const std::string& msg) {
        std::ostringstream ss;
        ss << "Function " << fn.name() << " returned an " << msg;
        return ss.str();
    });
}

```

```

    if (has_post_hooks) {
        // NOLINTNEXTLINE(bugprone-use-after-move)
        return call_post_hooks(fn, std::move(outputs), inputs);
    }
    return outputs;
}

```

可以看到，call_function()的核心逻辑就是执行fn()函数，这个fn函数指针是NodeTask的成员。而这个NodeTask是

```

    queue->push(
        NodeTask(graph_task, std::move(graph_root), std::move(input_buffer)));
struct NodeTask {
    std::weak_ptr<GraphTask> base_;
    std::shared_ptr<Node> fn_;
    // This buffer serves as an implicit "addition" node for all of the
    // gradients flowing here. Once all the dependencies are finished, we
    // use the contents of this buffer to run the function.
    InputBuffer inputs_;
    // When worker receives a task with isShutdownTask = true, it will immediately
    // exit. The engine sends a shutdown task to every queue upon its destruction.
    bool isShutdownTask_;

    int getReentrantDepth() const;

    NodeTask(
        // NOLINTNEXTLINE(modernize-pass-by-value)
        std::weak_ptr<GraphTask> base,
        std::shared_ptr<Node> fn,
        InputBuffer inputs,
        bool isShutdownTask = false)
        : base_(base),
          fn_(std::move(fn)),
          inputs_(std::move(inputs)),
          isShutdownTask_(isShutdownTask) {}
};

```

这样就知道所谓的NodeTask的成员fn_其实就是graph_root，而graph_root又是edge_list的第一项

```

auto graph_root = skip_dummy_node
    ? roots.at(0).function
    : std::make_shared<GraphRoot>(roots, inputs);

```

roots是一开始从Python调用C++函数的时候生成的，也就是在函数THPEngine_run_backward的实现里，相关的代码如

```

PyObject* THPEngine_run_backward(
    PyObject* self,
    PyObject* args,

```

```

    PyObject* kwargs) {
//...

    edge_list roots;
    roots.reserve(num_tensors);
    variable_list grads;
    grads.reserve(num_tensors);
    for (const auto i : c10::irange(num_tensors)) {
        PyObject* _tensor = PyTuple_GET_ITEM(tensors, i);
        THPUtils_assert(
            THPVariable_Check(_tensor),
            "element %d of tensors ",
            "tuple is not a Tensor",
            i);
        const auto& variable = THPVariable_Unpack(_tensor);
        TORCH_CHECK(
            !isBatchedTensor(variable),
            "torch.autograd.grad(outputs, inputs, grad_outputs) called inside ",
            "torch.vmap. We do not support the case where any outputs are ",
            "vmapped tensors (output ",
            i,
            " is being vmapped over). Please "
            "call autograd.grad() outside torch.vmap or file a bug report "
            "with your use case.")
        auto gradient_edge = torch::autograd::impl::gradient_edge(variable);
        THPUtils_assert(
            gradient_edge.function,
            "element %d of tensors does not require grad and does not have a grad_fn",
            i);
        roots.push_back(std::move(gradient_edge));

        //...
    }

    //...
}

```

gradient_edge的定义在torch/csrc/autograd/variable.cpp中:

```

Edge gradient_edge(const Variable& self) {
    // If grad_fn is null (as is the case for a leaf node), we instead
    // interpret the gradient function to be a gradient accumulator, which will
    // accumulate its inputs into the grad property of the variable. These
    // nodes get suppressed in some situations, see "suppress gradient
    // accumulation" below. Note that only variables which have `requires_grad =
    // True` can have gradient accumulators.
    if (const auto& gradient = self.grad_fn()) {

```



```

        return Edge(gradient, self.output_nr());
    } else {
        return Edge(grad_accumulator(self), 0);
    }
}

```

Edge的定义在torch/csrc/autograd/edge.h中，可以看出，Edge中的函数其实就是Variable中的grad_fn，而Variable

```

/// Represents a particular input of a function.
struct Edge {
    Edge() noexcept : function(nullptr), input_nr(0) {}

    Edge(std::shared_ptr<Node> function_, uint32_t input_nr_) noexcept
        : function(std::move(function_)), input_nr(input_nr_) {}

    /// Convenience method to test if an edge is valid.
    bool is_valid() const noexcept {
        return function != nullptr;
    }

    /// Required for use in associative containers.
    bool operator==(const Edge& other) const noexcept {
        return this->function == other.function && this->input_nr == other.input_nr;
    }

    bool operator!=(const Edge& other) const noexcept {
        return !(*this == other);
    }

    /// The function this `Edge` points to.
    std::shared_ptr<Node> function;

    /// The identifier of a particular input to the function.
    uint32_t input_nr;
};

```

参考

- PYTORCH 自动微分（二） <https://zhuanlan.zhihu.com/p/111874952>
- <https://zhuanlan.zhihu.com/p/69294347>
- <https://pytorch.org/blog/how-computational-graphs-are-executed-in-pytorch/>
- <https://www.cnblogs.com/rossiXYZ/p/15481235.html>